

xv6 study

(Lab: Xv6 and Unix utilities)

(0) xv6

Scheduling: Introduction

- ❑ xv6 is a reimplementation of UNIX for x86 and RISC-V
- ❑ It was created for educational purpose.
- ❑ List of System Call

<code>int fork()</code>	Create a process, return child's PID.
<code>int exit(int status)</code>	Terminate the current process; status reported to <code>wait()</code> . No return.
<code>int wait(int *status)</code>	Wait for a child to exit; exit status in <code>*status</code> ; returns child PID.
<code>int kill(int pid)</code>	Terminate process PID. Returns 0, or -1 for error.
<code>int getpid()</code>	Return the current process's PID.
<code>int sleep(int n)</code>	Pause for n clock ticks.
<code>int exec(char *file, char *argv[])</code>	Load a file and execute it with arguments; only returns if error.
<code>char *sbrk(int n)</code>	Grow process's memory by n bytes. Returns start of new memory.
<code>int open(char *file, int flags)</code>	Open a file; flags indicate read/write; returns an fd (file descriptor).
<code>int write(int fd, char *buf, int n)</code>	Write n bytes from buf to file descriptor fd; returns n.
<code>int read(int fd, char *buf, int n)</code>	Read n bytes into buf; returns number read; or 0 if end of file.
<code>int close(int fd)</code>	Release open file fd.
<code>int dup(int fd)</code>	Return a new file descriptor referring to the same file as fd.
<code>int pipe(int p[])</code>	Create a pipe, put read/write file descriptors in <code>p[0]</code> and <code>p[1]</code> .
<code>int chdir(char *dir)</code>	Change the current directory.
<code>int mkdir(char *dir)</code>	Create a new directory.
<code>int mknod(char *file, int, int)</code>	Create a device file.
<code>int fstat(int fd, struct stat *st)</code>	Place info about an open file into <code>*st</code> .
<code>int stat(char *file, struct stat *st)</code>	Place info about a named file into <code>*st</code> .
<code>int link(char *file1, char *file2)</code>	Create another name (file2) for the file file1.
<code>int unlink(char *file)</code>	Remove a file.

- ❑ “Lab: Xv6 and Unix utilities” helps us familiarize with xv6 and it's system calls.

- ◆ There are five problems : sleep, pingpong, primes, find and xargs.

(1) sleep

sleep : solution

```
$ make qemu
...
init: starting sh
$ sleep 10
(nothing happens for a little while)
$
```

```
int main(int argc, char *argv[]) {
    if (argc != 2) {
        printf("Usage: sleep <time>\n");
        exit(1);
    }

    int time = atoi(argv[1]);

    if (time <= 0) {
        printf("Error: sleep time must be a positive integer\n");
        exit(1);
    }

    sleep(time);

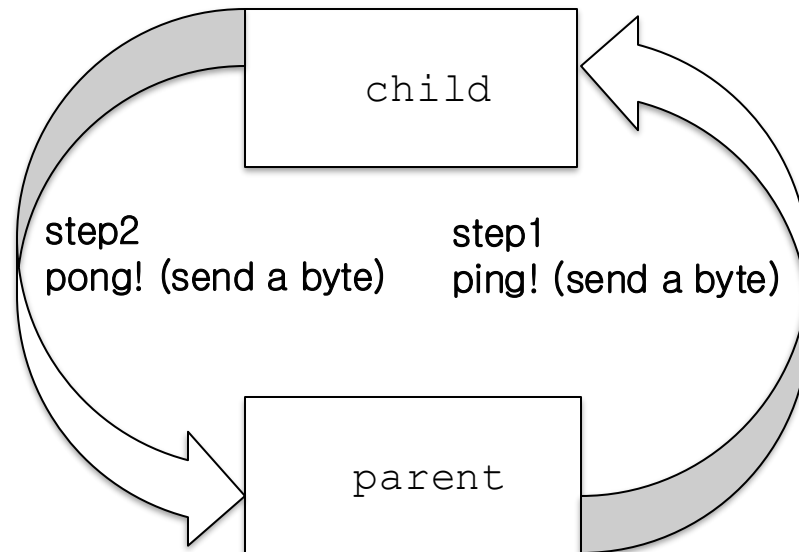
    exit(0);
}
```

(2) pingpong


pingpong : solution

```
int p[2];
int q[2];
pipe(p);
pipe(q);
char c;
if(fork() == 0) //child
{
    read(p[0], &c, 1);
    printf("%d: received ping\n", getpid());
    write(q[1], &c, 1);
}
else //parent
{
    write(p[1], &c, 1);
    read(q[0], &c, 1);
    printf("%d: received pong\n", getpid());
}
exit(0);
```

- Step1
 - ◆ The parent sends a byte to the child
 - ◆ The child receive a byte
 - print "<pid>:received ping
- Step2
 - ◆ The child sends a byte to the parent
 - ◆ The parent receive a byte
 - print "<pid>:received pong



pipe, read, write



▣ Step1

- ♦ call the system-call "pipe" with an integer array of two elements.
- ♦ The first element of the array hold the file descriptor for reading.
- ♦ The second element of the array hold the file descriptor for writing.

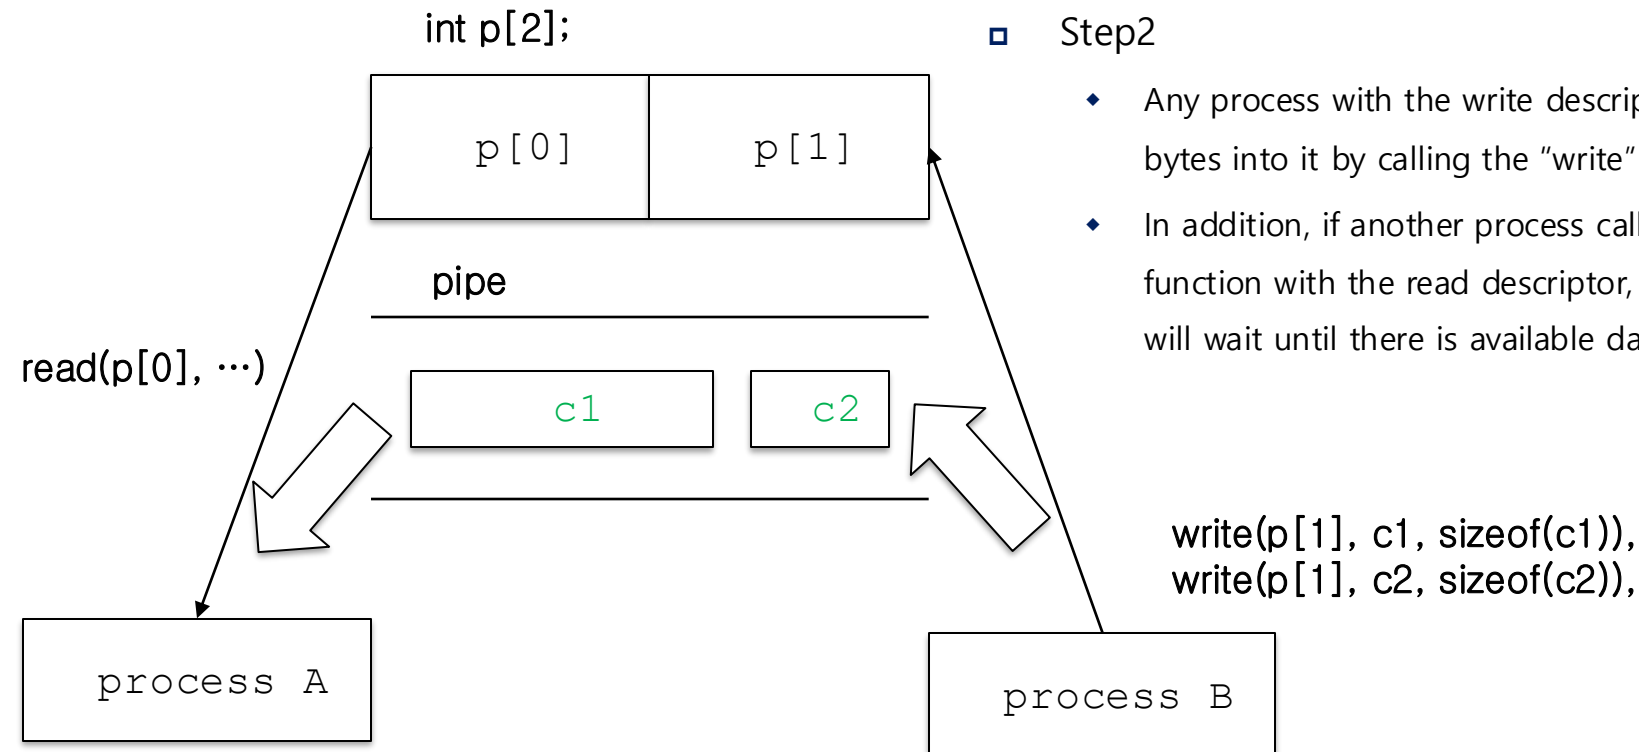
pipe, read, write

Step1

- ♦ call the system-call "pipe" with an integer array of two elements.
- ♦ The first element of the array hold the file descriptor for reading.
- ♦ The second element of the array hold the file descriptor for writing.

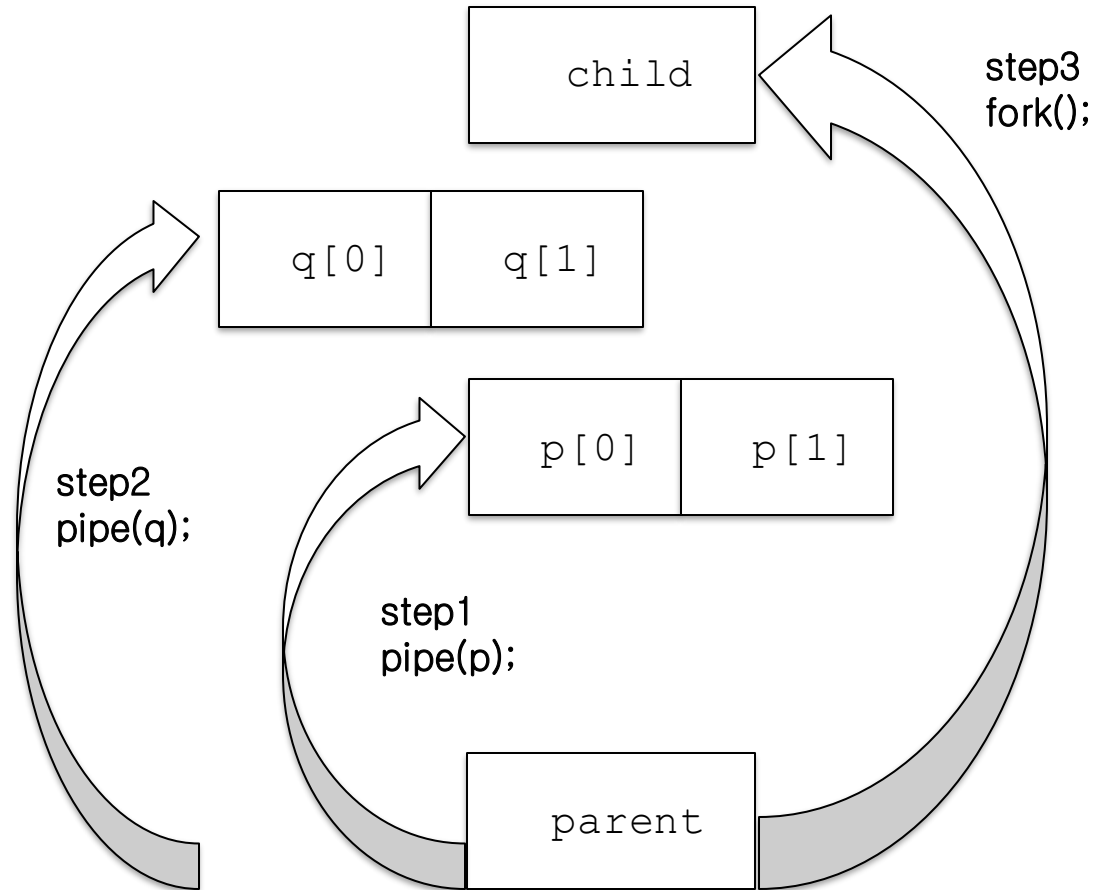
Step2

- ♦ Any process with the write descriptor can write bytes into it by calling the "write" function.
- ♦ In addition, if another process call the "read" function with the read descriptor, the process will wait until there is available data in pipe




pingpong : solution

```
int p[2];
int q[2];
pipe(p);
pipe(q);
char c;
if(fork() == 0) //child
{
    read(p[0], &c, 1);
    printf("%d: received ping\n", getpid());
    write(q[1], &c, 1);
}
else //parent
{
    write(p[1], &c, 1);
    read(q[0], &c, 1);
    printf("%d: received pong\n", getpid());
}
exit(0);
```




pingpong : solution

```
int p[2];
int q[2];
pipe(p);
pipe(q);
char c;
if(fork() == 0) //child
{
    read(p[0], &c, 1);
    printf("%d: received ping\n", getpid());
    write(q[1], &c, 1);
}
else //parent
{
    write(p[1], &c, 1);
    read(q[0], &c, 1);
    printf("%d: received pong\n", getpid());
}
exit(0);
```



pingpong : solution

```
int p[2];
int q[2];
pipe(p);
pipe(q);
char c;
if(fork() == 0) //child
{
    read(p[0], &c, 1);
    printf("%d: received ping\n", getpid());
    write(q[1], &c, 1);
}
else //parent
{
    write(p[1], &c, 1);
    read(q[0], &c, 1);
    printf("%d: received pong\n", getpid());
}
exit(0);
```



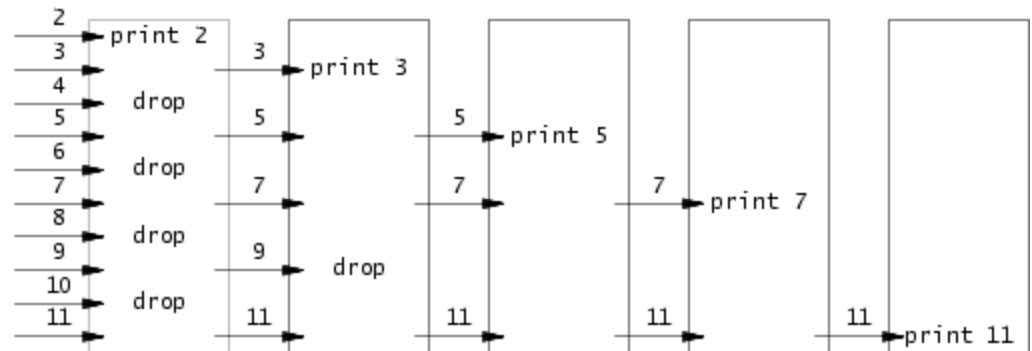
(3) primes

primes : solution

```
int ps[2];
pipe(ps);
{//step1
    ...
}

int prime;
while(read(ps[0], &prime, sizeof(int))){
    printf("prime %d\n", prime);
    int new_ps[2];
    pipe(new_ps);
    if(!fork())
        {//step2
            //
        }
    else
    {
        wait(0);
    }
    close(ps[0]);
    close(new_ps[1]);
    ps[0] = new_ps[0];
}
```

- step1
 - ◆ The first process feed the number 2 through 280 into the left end of the pipeline
- step2
 - ◆ For each prime p, eliminates the multiples of p, and feed the remained numbers into the left end of the pipeline




primes : solution

```
int ps[2];
pipe(ps);
{//step1
    for(int i=2;i<=280;i++)write(ps[1], &i, sizeof(i));
}

int prime;
while(read(ps[0], &prime, sizeof(int))){
    printf("prime %d\n", prime);
    int new_ps[2];
    pipe(new_ps);
    if(!fork())
        {//step2
            //
        }
    else
    {
        wait(0);
    }
    close(ps[0]);
    close(new_ps[1]);
    ps[0] = new_ps[0];
}
```

□ step1

- ◆ The first process feed the number 2 through 280 into the left end of the pipeline




primes : solution

```
...
int prime;
while(read(ps[0], &prime, sizeof(int))){
    printf("prime %d\n", prime);
    int new_ps[2];
    pipe(new_ps);
    if(!fork())
        {//step2
            //
        }
    else
    {
        wait(0);
    }
    close(ps[0]);
    close(new_ps[1]);
    ps[0] = new_ps[0];
}
```

□ step2

- ◆ For each prime p , eliminates the multiples of p , and feed the remained numbers into the left end of the pipeline




primes : solution

```
...
while(read(ps[0], &prime, sizeof(int))){
    ....
    int new_ps[2];
    pipe(new_ps);
    if(!fork())
    {//step2
        int num;
        while(read(ps[0], &num, sizeof(int)))
            if(num % prime != 0) write(new_ps[1], &num, sizeof(int));
        exit(0);
    }
    else
    {
        wait(0);
    }
    close(ps[0]);
    close(new_ps[1]);
    ps[0] = new_ps[0];
}
}
```

□ step2

- ◆ For each prime p , eliminates the multiples of p , and feed the remained numbers into the left end of the pipeline

parent




primes : solution

```
...
while(read(ps[0], &prime, sizeof(int))){
    ....
    int new_ps[2];
    pipe(new_ps);
    if(!fork())
    {//step2
        int num;
        while(read(ps[0], &num, sizeof(int)))
            if(num % prime != 0) write(new_ps[1], &num, sizeof(int));
        exit(0);
    }
    else
    {
        wait(0);
    }
    close(ps[0]);
    close(new_ps[1]);
    ps[0] = new_ps[0];
}
}
```

□ step2

- ◆ For each prime p , eliminates the multiples of p , and feed the remained numbers into the left end of the pipeline

parent




primes : solution

```
...
while(read(ps[0], &prime, sizeof(int))){
    ....
    int new_ps[2];
    pipe(new_ps);
    if(!fork())
    {//step2
        int num;
        while(read(ps[0], &num, sizeof(int)))
            if(num % prime != 0) write(new_ps[1], &num, sizeof(int));
        exit(0);
    }
    else
    {
        wait(0);
    }
    close(ps[0]);
    close(new_ps[1]);
    ps[0] = new_ps[0];
}
}
```

□ step2

- ◆ For each prime p , eliminates the multiples of p , and feed the remained numbers into the left end of the pipeline



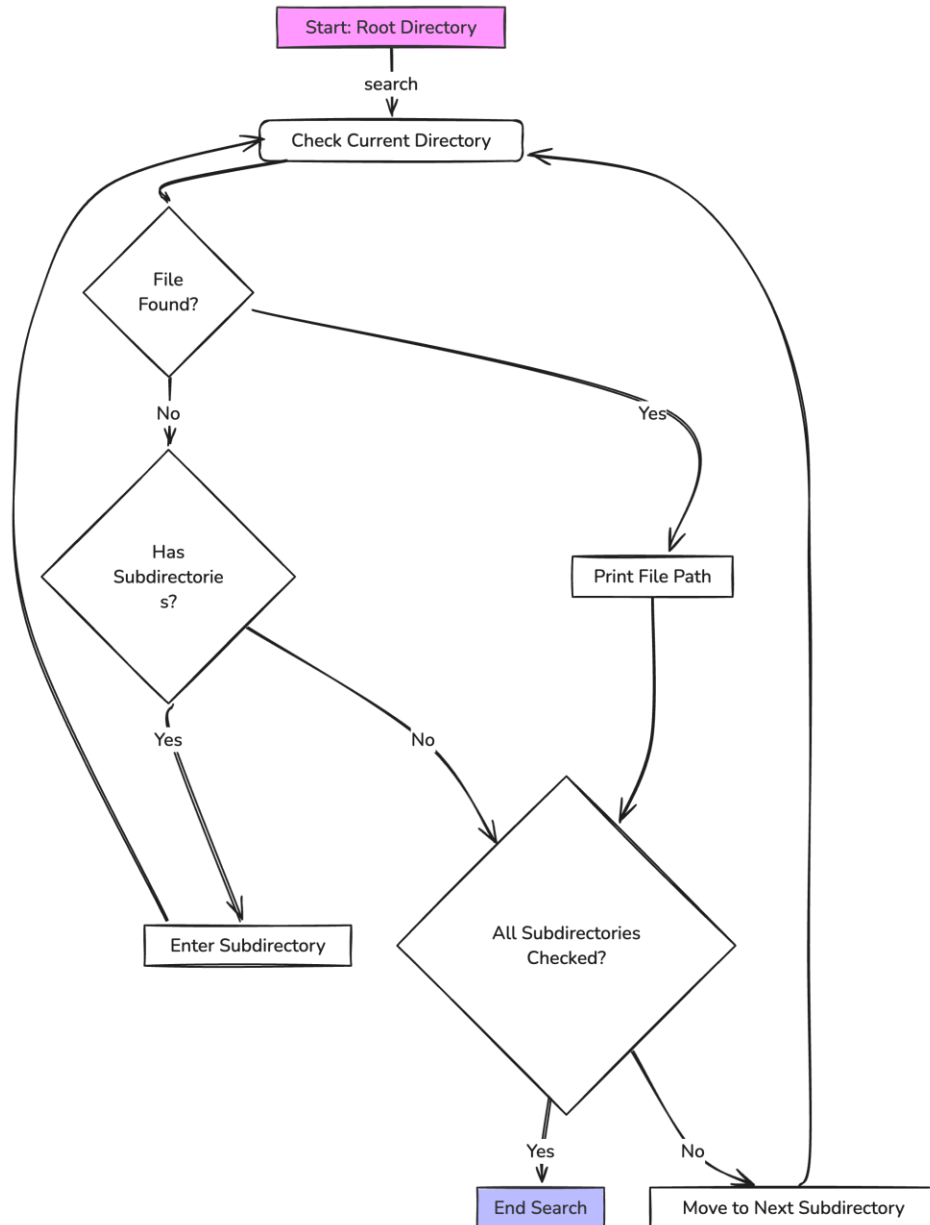
$ps[0] = \{ \textcolor{red}{2}, 3, \textcolor{red}{4}, 5, \textcolor{red}{6}, \dots, 279, \textcolor{red}{280} \}$

(4) find

find : overview

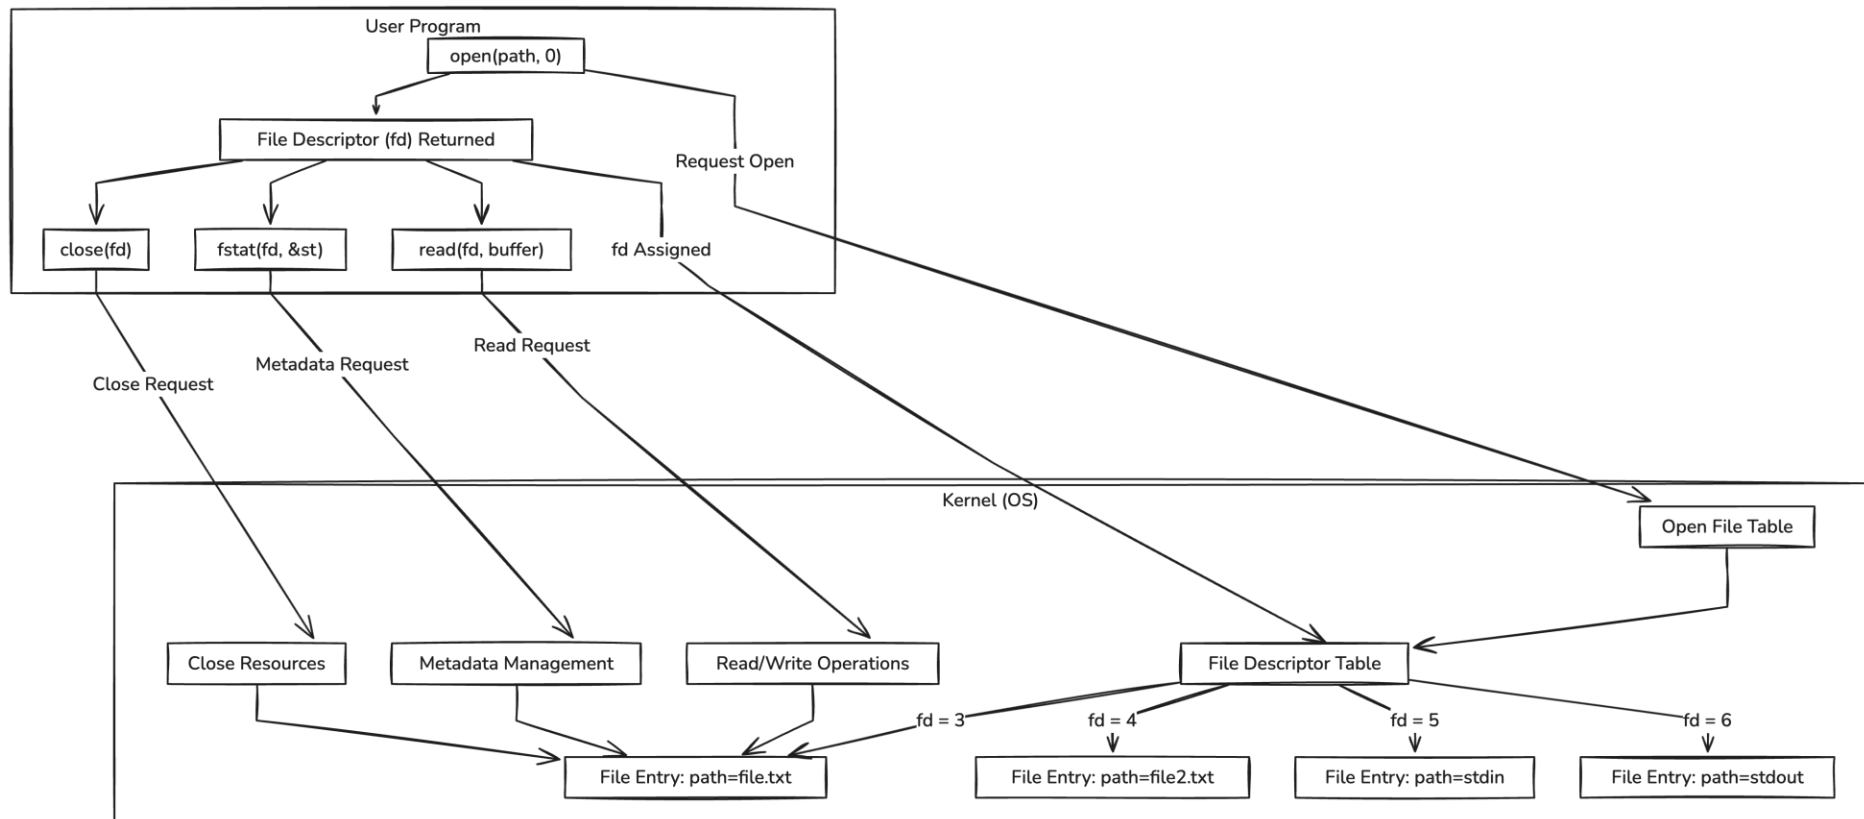
- Search for files recursively in a directory.

User Program (find)		Kernel (OS)
open(path, 0)	--->	Open the directory file Return File Descriptor FD
	<---	
fstat(fd, &st)	--->	Fetch metadata (st.type) Return file status
	<---	
read(fd, &de, sizeof(de))	--->	Read directory entry Return entry information
	<---	
stat(buf, &st)	--->	Get metadata for file path Return file status
	<---	
close(fd)	--->	Close the file descriptor Release resources
	<---	



find : overview

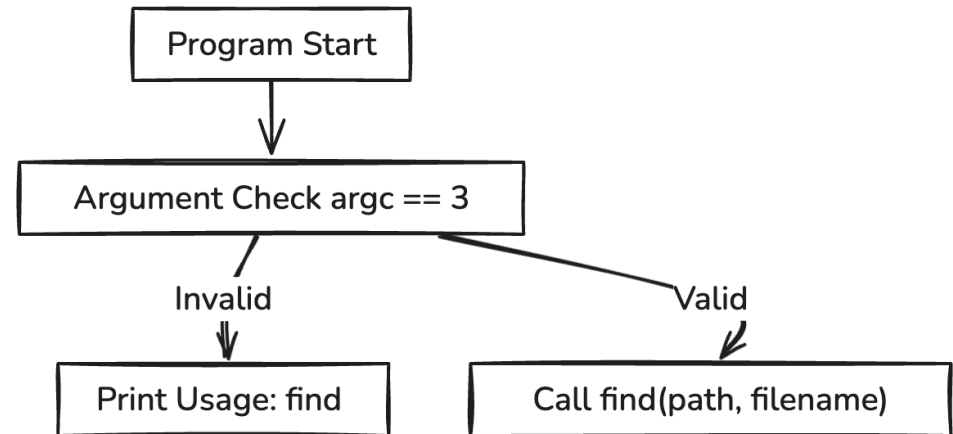
System Call Flow



find : solution

▣ Main Function

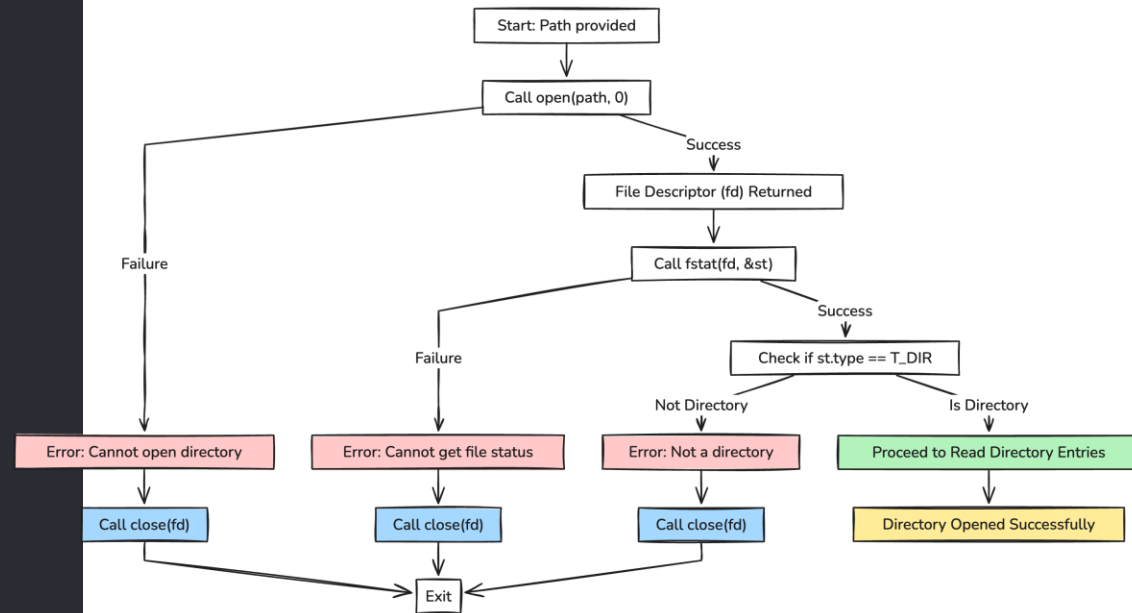
```
73 int main(int argc, char *argv[]) {  
74     if (argc != 3) {  
75         printf("Usage: find <path> <filename>\n");  
76         exit(1);  
77     }  
78     find(argv[1], argv[2]);  
79     exit(0);  
80 }  
81  
82
```



find : solution

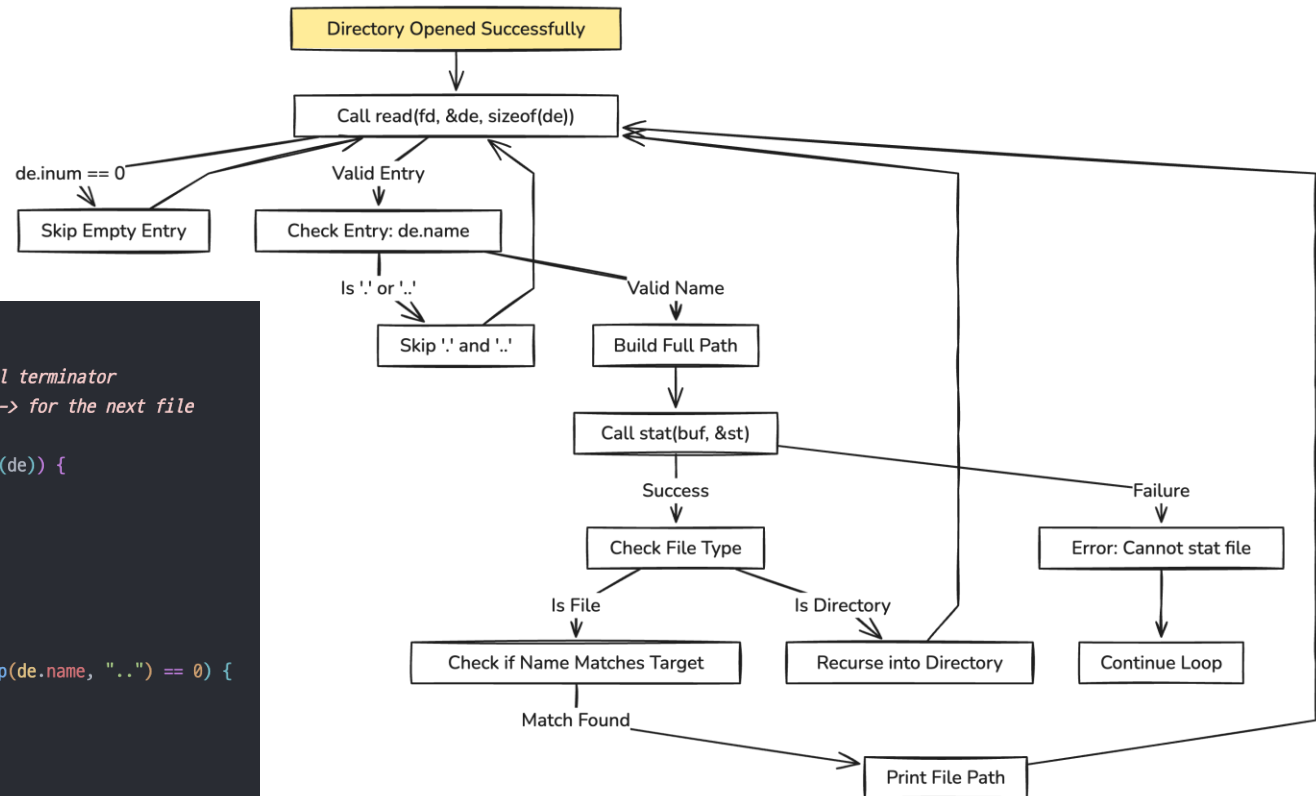
- Find Function : Opening and Verifying the Directory: `open()` and `fstat()`

```
6 void find(const char *path, const char *target) {
7     char buf[512], *p;
8     int fd;
9     struct dirent de;
10    struct stat st;
11
12    // open the directory
13    if ((fd = open(path, 0)) < 0) {
14        fprintf(2, "find: cannot open %s\n", path);
15        return;
16    }
17
18    // get the status of the directory
19    if (fstat(fd, &st) < 0) {
20        fprintf(2, "find: cannot open %s\n", path);
21        close(fd);
22        return;
23    }
24
25    // ensure it is a directory
26    if (st.type != T_DIR) {
27        fprintf(2, "find: %s is not a directory\n", path);
28        close(fd);
29        return;
30    }
```



find : solution

- Find Function : Reading Directory Entries: `read()` and Iterating Entries

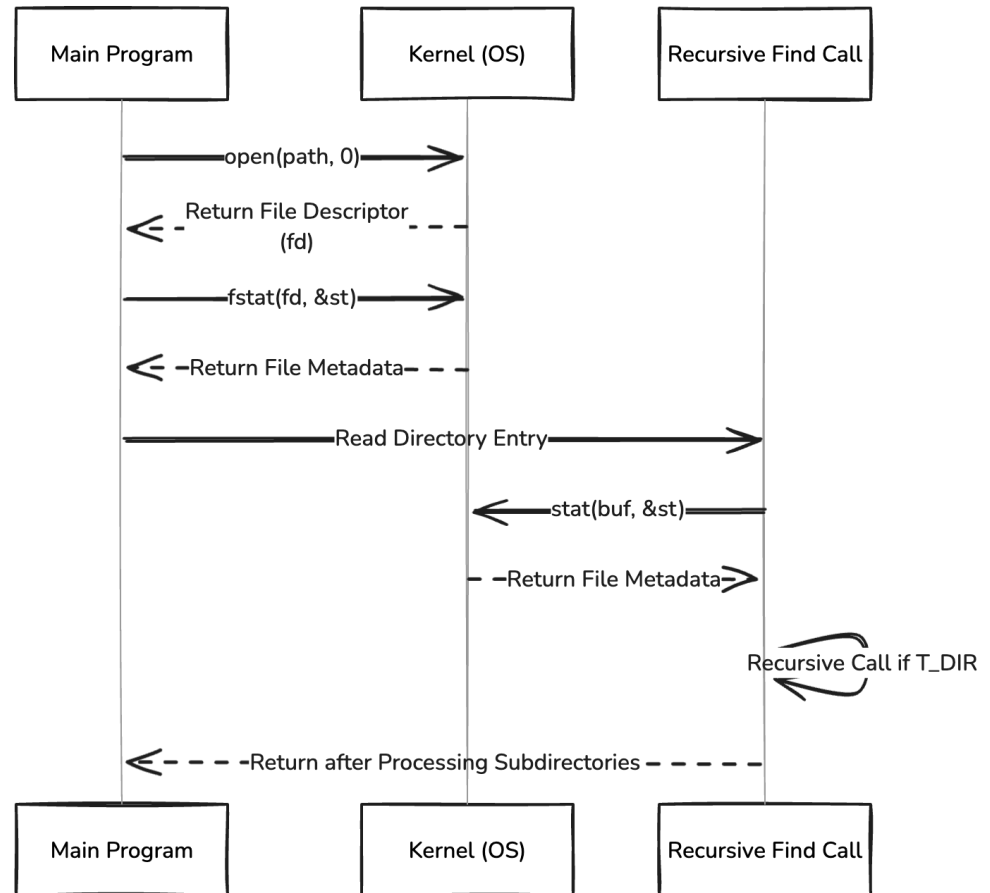


```
32 // read directory entries
33 strcpy(buf, path);
34 p = buf+strlen(buf); // point p to the null terminator
35 *p++ = '/'; // append a slash to the path -> for the next file
36
37 while (read(fd, &de, sizeof(de)) == sizeof(de)) {
38     // skip empty directory entries
39     if (de.inum == 0) {
40         continue;
41     }
42
43     // skip "." and ".."
44     if (strcmp(de.name, ".") == 0 || strcmp(de.name, "..") == 0) {
45         continue;
46     }
47
48     // Build the full path
49     memmove(p, de.name, DIRSIZ);
50     // ensure the path is null-terminated
51     p[DIRSIZ] = 0;
52
53     // get the status of the current file
54     if (stat(buf, &st) < 0) {
55         printf("find: cannot stat %s\n", buf);
56         continue;
57     }
```

find : solution

Find Function : Path Construction and Recursive Call

```
59 // check if the current file matches the target
60 if (strcmp(de.name, target) == 0) {
61     printf("%s\n", buf);
62 }
63
64 // if it's a directory, recurse into it
65 if (st.type == T_DIR) {
66     find(buf, target);
67 }
68 }
69
70 close(fd);
```

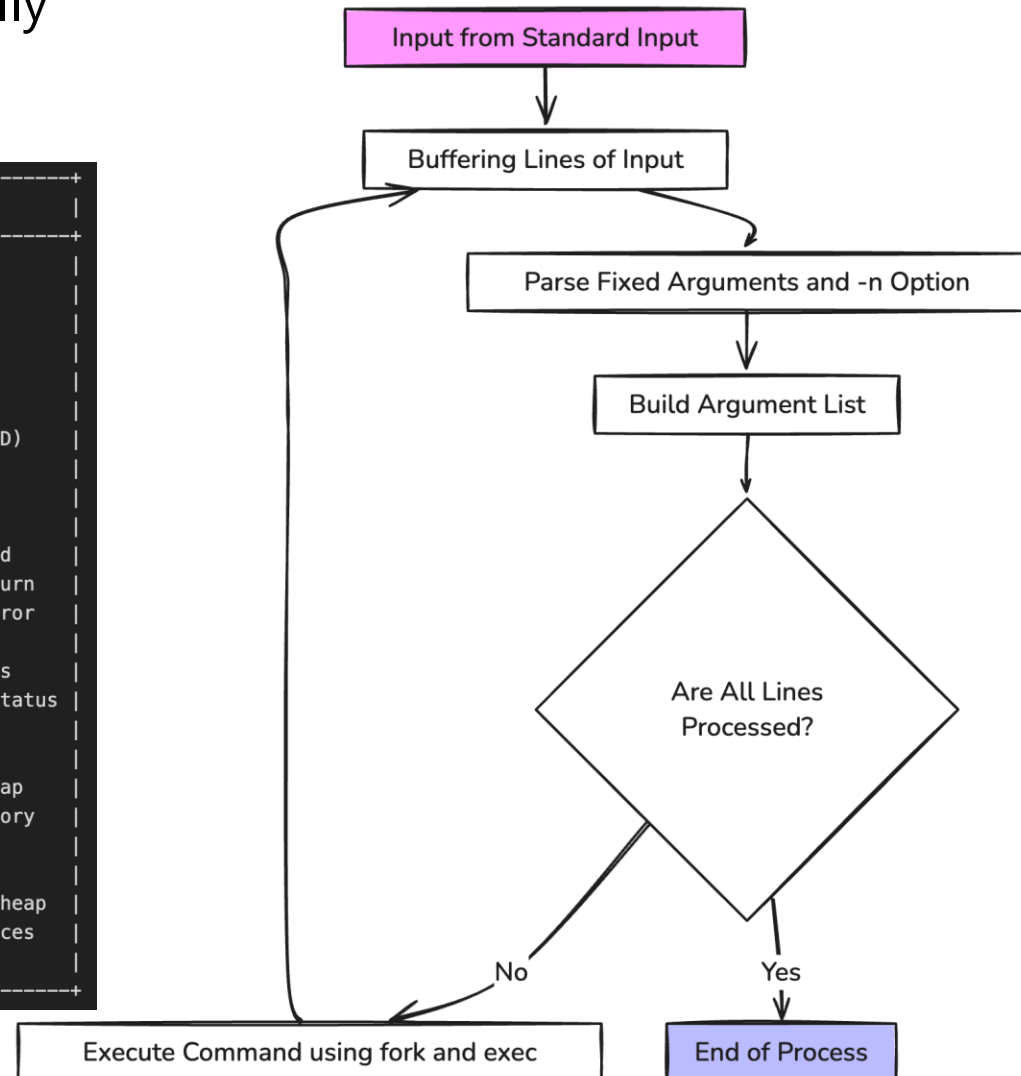


(5) xargs

xargs : overview

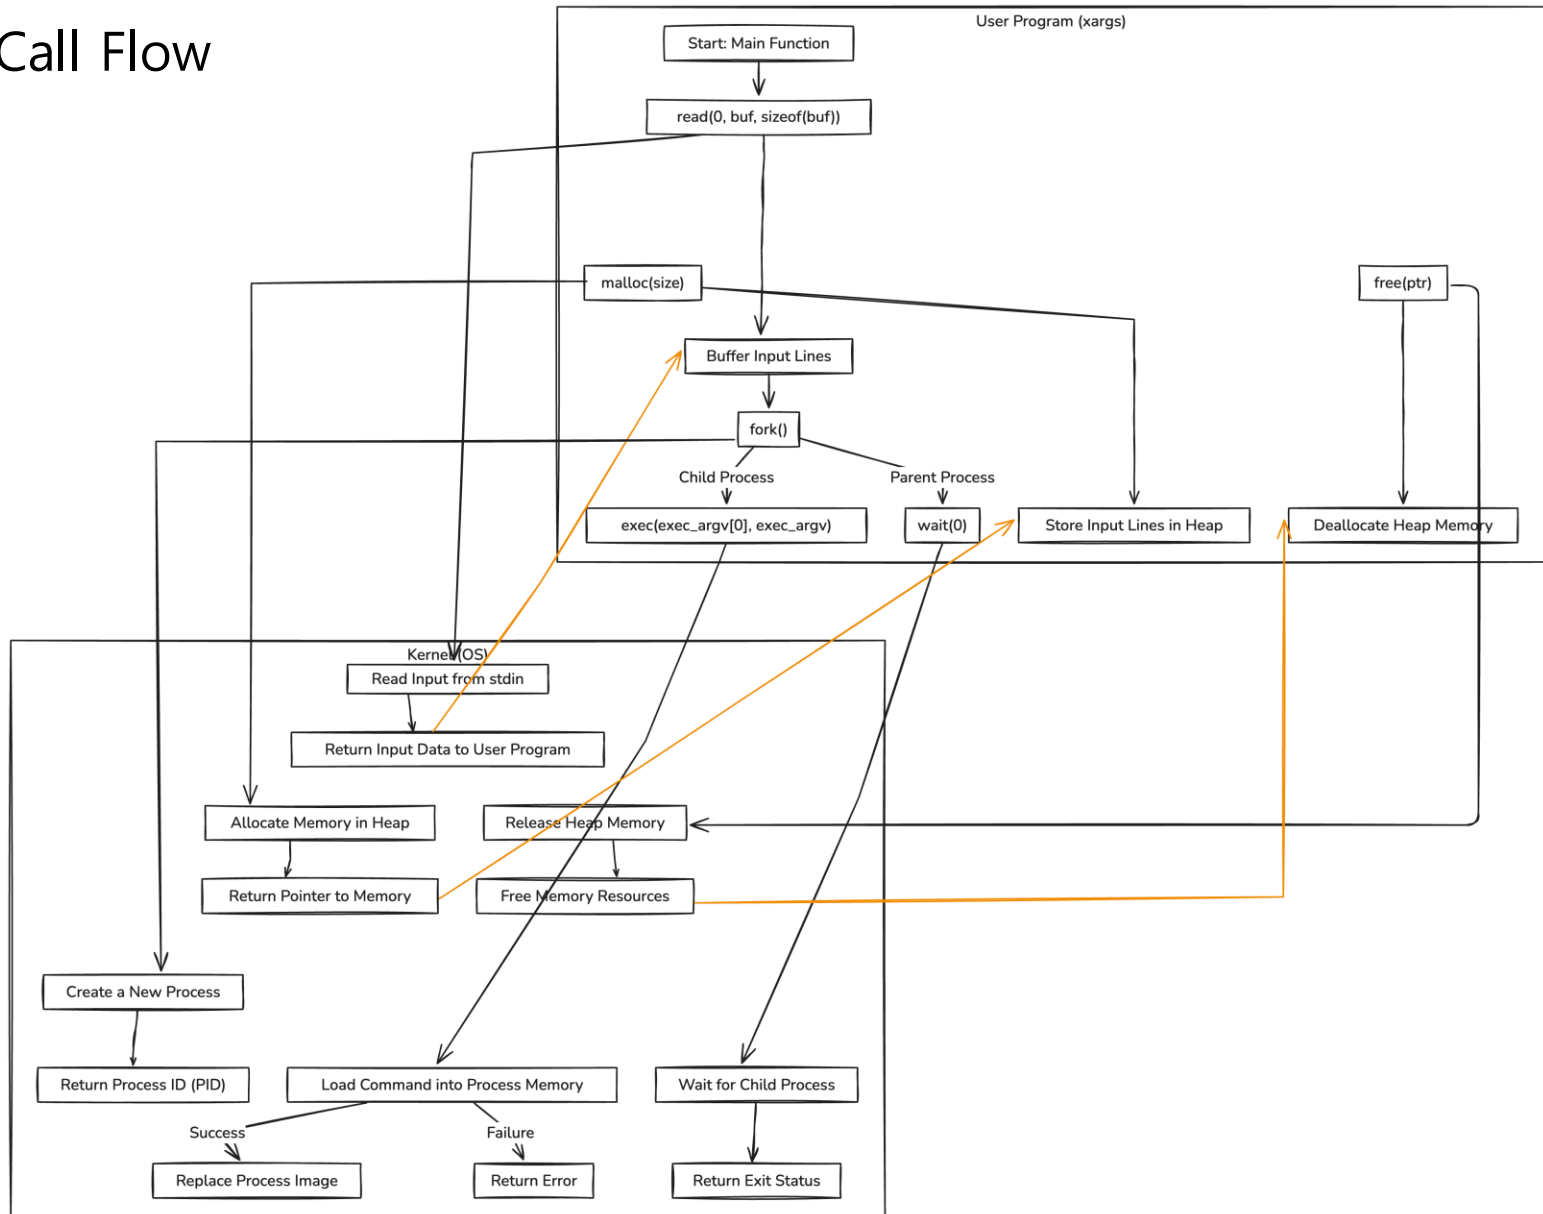
- Execute commands dynamically using input lines.

User Program (xargs)		Kernel (OS)
<code>read(0, buf, sizeof(buf))</code>	---->	Read input from stdin Return buffer data
	<---	
<code>fork()</code>	---->	Create a new process Return process ID (PID)
	<---	
<code>exec(exec_argv[0], exec_argv)</code>	---->	Replace process image Load specified command If successful, no return On failure, return error
	<---	
<code>wait(0)</code>	---->	Wait for child process Return child's exit status
	<---	
<code>malloc(size)</code>	---->	Allocate memory in heap Return pointer to memory
	<---	
<code>free(ptr)</code>	---->	Deallocate memory in heap Release memory resources
	<---	



xargs : overview

■ System Call Flow

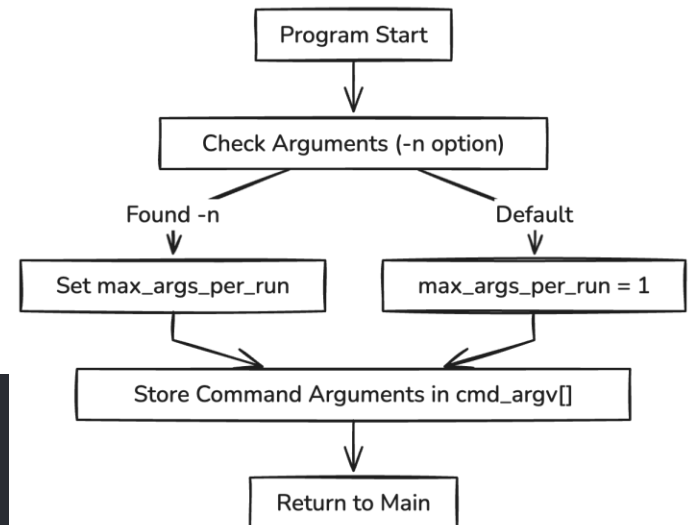


xargs : solution

Argument Parsing

```
113 int main(int argc, char **argv) {
114     if (argc < 2) {
115         printf("Usage: xargs [-n num] <command> [args...]\n");
116         exit(1);
117     }
118
119     int max_args_per_run = 1, start_index = 1, cmd_argc = 0;
120     char *cmd_argv[MAXARG];
121     parse_arguments(argc, argv, &max_args_per_run, &start_index, cmd_argv, &cmd_argc);
```

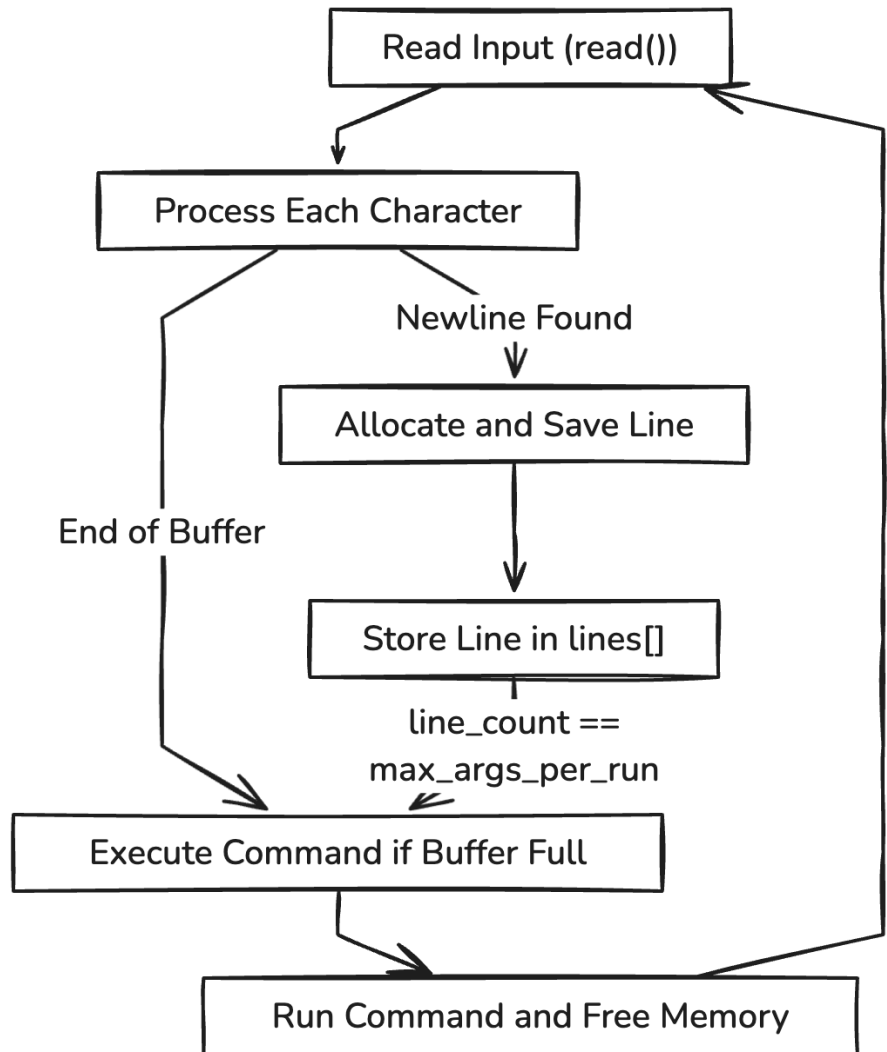
```
47 void parse_arguments(int argc, char **argv, int *max_args_per_run, int *start_index, char **cmd_argv, int *cmd_argc) {
48     if (argc > 2 && strcmp(argv[1], "-n") == 0) {
49         *max_args_per_run = atoi(argv[2]);
50         if (*max_args_per_run < 1) {
51             printf("xargs: invalid value for -n\n");
52             exit(1);
53         }
54         *start_index = 3;
55     }
56     for (int i = *start_index; i < argc && *cmd_argc < MAXARG - 1; i++) {
57         cmd_argv[(*cmd_argc)++] = argv[i];
58     }
59     cmd_argv[*cmd_argc] = 0;
60 }
```



xargs : solution

■ Reading Input and Storing Lines


```
86 void process_input(char **cmd_argv, int cmd_argc, int max_args_per_run) {
87     char buf[BUFFER_SIZE], line[BUFFER_SIZE];
88     int n, line_idx = 0, line_count = 0;
89     char *lines[MAXARG];
90
91     while ((n = read(0, buf, sizeof(buf))) > 0) {
92         for (int i = 0; i < n; i++) {
93             if (buf[i] == '\n') {
94                 line[line_idx] = 0;
95                 allocate_and_save_line(line, lines, &line_count);
96                 if (line_count == max_args_per_run) {
97                     execute_lines(cmd_argv, cmd_argc, lines, line_count);
98                     line_count = 0;
99                 }
100                 line_idx = 0;
101             } else {
102                 line[line_idx++] = buf[i];
103             }
104         }
105     }
106     if (line_idx > 0) {
107         line[line_idx] = 0;
108         allocate_and_save_line(line, lines, &line_count);
109     }
110     execute_lines(cmd_argv, cmd_argc, lines, line_count);
111 }
```




xargs : solution

▣ Memory Allocation and Cleanup

```
62 void allocate_and_save_line(char *line, char **lines, int *line_count) {
63     char *saved_line = malloc(strlen(line) + 1);
64     if (!saved_line) {
65         printf("xargs: malloc failed\n");
66         exit(1);
67     }
68     strcpy(saved_line, line);
69     remove_quotes(saved_line);
70     lines[(*line_count)++] = saved_line;
71 }
```



```
80 void cleanup_lines(char **lines, int line_count) {
81     for (int i = 0; i < line_count; i++) {
82         free(lines[i]);
83     }
84 }
```



xargs : solution

■ Executing Commands

```
25 void run_command(int cmd_argc, char **cmd_argv, int line_count, char **lines) {
26     char *exec_argv[MAXARG];
27     int i, j = 0;
28
29     for (i = 0; i < cmd_argc; i++) {
30         exec_argv[j++] = cmd_argv[i];
31     }
32     for (i = 0; i < line_count; i++) {
33         exec_argv[j++] = lines[i];
34     }
35
36     exec_argv[j] = 0;
37
38     if (fork() == 0) {
39         exec(exec_argv[0], exec_argv);
40         printf("xargs: exec failed for command %s\n", exec_argv[0]);
41         exit(1);
42     } else {
43         wait(0);
44     }
45 }
```

