

Traps and System Calls

Trap

- System call
 - e.g. ecall
 - Exception
 - e.g. illegal instruction
 - Kernel kills the offending program
 - Device interrupt
 - e.g. disk finishes read/write
-
1. Trap forces transfer of control into the kernel
 2. Kernel executes appropriate handler code (save then restore previous state)
 3. Return to user code where original code resumes (without needing to know anything special happened)

RISC-V Trap Machinery

- stvec : address of the trap handler
- sepc: program counter saved (pc is overwritten with value in stvec)
- scause: trap reason (number)
- sscratch: avoid overwriting user registers
- sstatus: control whether device interrupts are enabled

Traps from User Space

uservec → usertrap → ..trap handling.. → usertrapret → userret

- Trampoline page : uservec
- Uservec: save 32 register values (using sscratch register)
- Usertrap: determine case of the trap, process, and return
- Usertrapret: prepare for future trap from user space

Page Fault Exceptions

Exception handling in xv6:

- User space: kills the faulting process
- Kernel: kernel panics

Page Fault:

- Load page fault
- Store page fault
- Instruction page fault

RISC-V Assembly

```

int g(int x) {
    return x+3;
}

int f(int x) {
    return g(x);
}

void main(void) {
    printf("%d %d\n", f(8)+1, 13);
    exit(0);
}

```

```

0000000000000000 <g>:
int g(int x) {
    0:      1141                addi    sp,sp,-16
    2:      e406                sd      ra,8(sp)
    4:      e022                sd      s0,0(sp)
    6:      0800                addi    s0,sp,16
    return x+3;
}
    8:      250d                addiw   a0,a0,3
    a:      60a2                ld      ra,8(sp)
    c:      6402                ld      s0,0(sp)
    e:      0141                addi    sp,sp,16
    10:     8082                ret

0000000000000012 <f>:
int f(int x) {
    12:     1141                addi    sp,sp,-16
    14:     e406                sd      ra,8(sp)
    16:     e022                sd      s0,0(sp)
    18:     0800                addi    s0,sp,16
    return g(x);
}
    1a:     250d                addiw   a0,a0,3
    1c:     60a2                ld      ra,8(sp)
    1e:     6402                ld      s0,0(sp)
    20:     0141                addi    sp,sp,16
    22:     8082                ret

```

```

0000000000000024 <main>:
void main(void) {
    24:      1141          addi    sp,sp,-16
    26:      e406          sd      ra,8(sp)
    28:      e022          sd      s0,0(sp)
    2a:      0800          addi    s0,sp,16
    printf("%d %d\n", f(8)+1, 13);
    2c:      4635          li      a2,13
    2e:      45b1          li      a1,12
    30:      00001517      auipc   a0,0x1
    34:      86050513      addi    a0,a0,-1952 # 890
    <malloc+0xf6>
    38:      6aa000ef      jal     6e2 <printf>
    exit(0);
    3c:      4501          li      a0,0
    3e:      2a2000ef      jal     2e0 <exit>
**

```

1. Which registers contain arguments to functions?
For example, which register holds 13 in main's call to printf?

-> Registers a0-a7.

2. Where is the call to function f in the assembly code for main? Where is the call to g? (Hint: the compiler may inline functions.)

-> No explicit calls are made.

3. At what address is the function printf located?

-> 0x6e2

4. What value is in the register ra just after the jalr to printf in main?

-> address 0x3c -- the address of the next instruction.


```
unsigned int i = 0x00646c72;  
printf("H%x Wo%s", 57616, (char *) &i);
```

5. What is the output of the given code?

-> HE110 World

57616 in hex is 0xE110

0x00646c72 in memory is stored as

72:(ASCII) r, 6c:l, 64:d, 00:\0 (ordered by lower to higher address)

For big endian, i has to be set differently.

i = 0x726c6400 would be

00:\0, 64:d, 6c:l, 72:r, and big endian will store bytes from higher address.

```
printf("x=%d y=%d", 3);
```

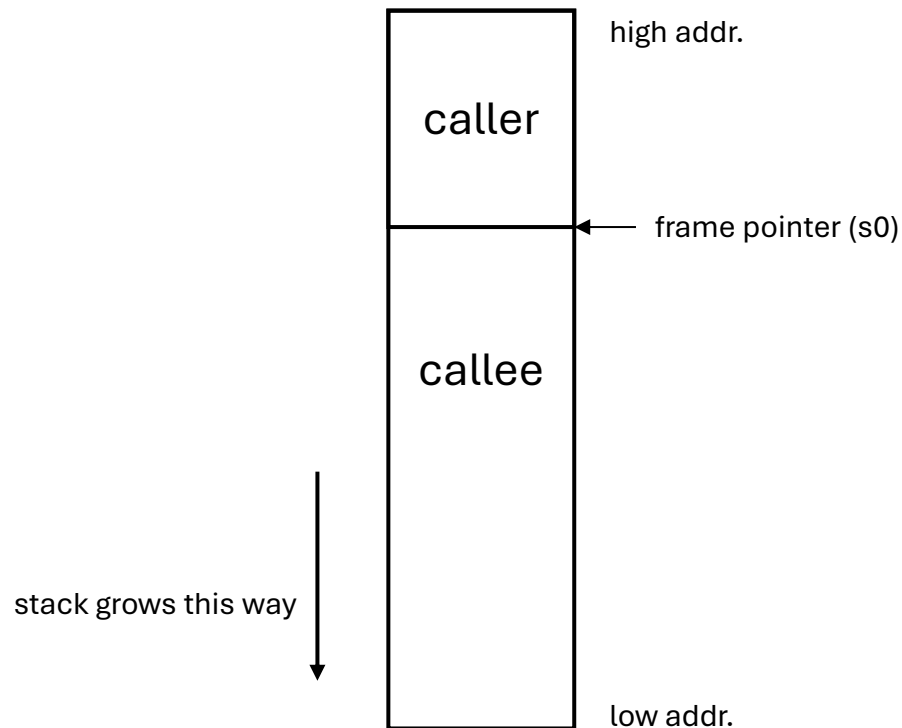
6. What is going to be printed after 'y='? Why does this happen?

-> some garbage value

Backtrace

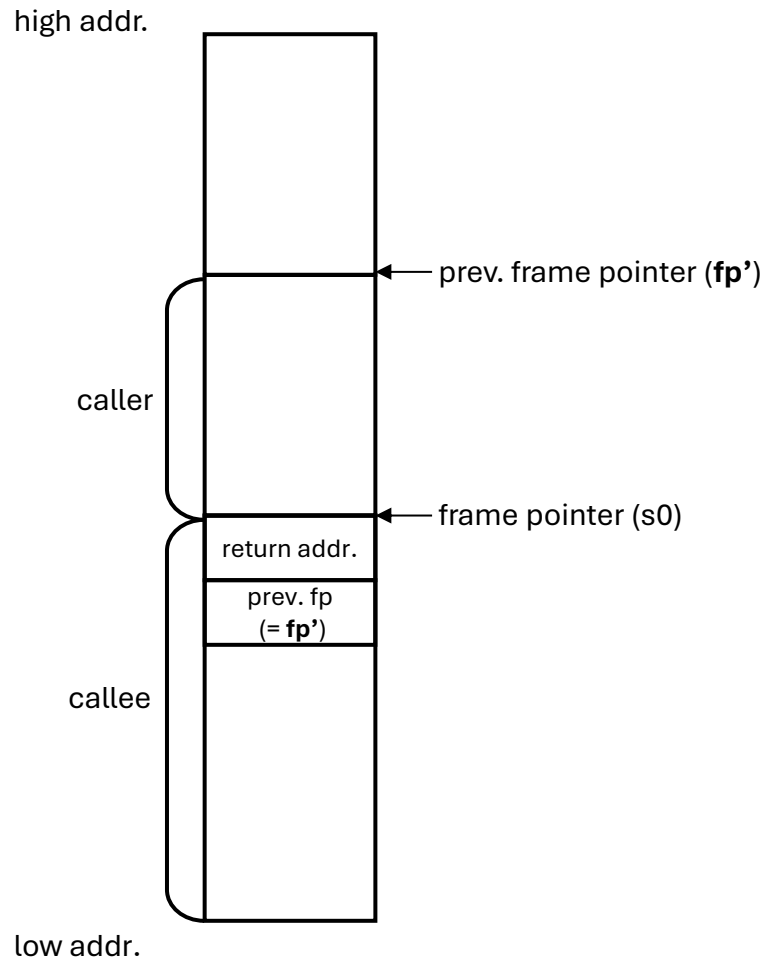
Reading from registers – read frame pointer

- s0 contains the current frame pointer (memory address)
- frame pointer points to the bottom of the callee stack



```
28 + // current value of the fp
29 + static inline uint64
30 + r_fp()
31 + {
32 +     uint64 x;
33 +     asm volatile("mv %0, s0" : "=r"(x));
34 +     return x;
35 + }
36 +
```

Restoring Return Address and prev. FP



```
62 + static inline uint64 get_return_addr(uint64 fp)
63 + {
64 +     return *((uint64*)(fp - 8));
65 + }
66 +
67 + static inline uint64 get_saved_fp(uint64 fp)
68 + {
69 +     return *((uint64*)(fp - 16));
70 + }
71 +
```

Walk up the stack until fp is out of the current page

```
72 + void backtrace()
73 + {
74 +     uint64 fp = r_fp();           // fp when backtrace() called
75 +     const uint64 pg = PGROUNDDOWN(fp); // page of the kernel stack
76 +     printf("backtrace:\n");
77 +     do {
78 +         printf("%p\n", (void*)get_return_addr(fp));
79 +         fp = get_saved_fp(fp); // restore saved fp
80 +     } while (pg == PGROUNDDOWN(fp));
81 + }
82 +
```

Alarm – Yi, Cherry

Requirements

- System Call Implementation

- Sigalarm(ticks, handler)

- Kernel should trigger the handler ftn every ticks
 - If sigalarm(0,0) is called, periodic alarms should stop

- Sigreturn()

- Handler ftn must call sigreturn to restore the process's state after execution

```
34
35 void
36 periodic()
37 {
38     count = count + 1;
39     printf("alarm!\n");
40     sigreturn();
41 }
42
```

```
146 void
147 slow_handler()
148 {
149     count++;
150     printf("alarm!\n");
151     if (count > 1) {
152         printf("test2 failed: alarm handler called more than once\n");
153         exit(1);
154     }
155     for (int i = 0; i < 1000*500000; i++) {
156         asm volatile("nop"); // avoid compiler optimizing away loop
157     }
158     sigalarm(0, 0);
159     sigreturn();
160 }
161
```

```
165 void
166 dummy_handler()
167 {
168     sigalarm(0, 0);
169     sigreturn();
170 }
171
```

User Program	Kernel (Trap Handling)
-- [System Call] -->	--> Handle syscall (e.g., sys_sigalarm)
<---- [Return] -----	
-- [Execution] --> Timer Interrupt	
----- [Trap to Kernel] -----	
	usertrap() checks which interrupt occurred
	If timer and alarm condition met:
	- Save state
	- Set `epc` to alarm handler
--- [Return to Alarm Handler] --> User alarm handler executes	
-- [sigreturn System Call] -->	
----- [Trap to Kernel] -----	
	- Restore saved state
	- Resume original code execution

Test case & Key points to implement correctly

- Test0: Basic Alarm Handler Execution
- Test1: Multiple Alarm Handler Calls & Correct State Restoration
- Test2: Reentrant Alarm Prevention
 - `slow_handler()`는 오래 실행되며, 실행 중에 두 번째로 호출되면 안 됨.
- Test3: Register `a0` Preservation
 - `sigreturn()`이 호출될 때 레지스터 값이 변조되지 않으면 테스트 통과.

Test case & Key points to implement correctly

- **Timer Interrupt** : check if it is
- **State Management** : save the process state / restore when handler finishes
- **Re-entrant Handler Prevention** : use flag

Modified files & function

- Proc.c - sigalarm, sigreturn
 - 프로세스의 주요 동작 및 상태 관리 관련 기능을 제공.
- Sysproc.c - sys_sigalarm, sys_sigreturn
 - 사용자 프로그램에서 호출하는 시스템 콜 처리.
- Trap.c - usertrap
 - 예외, 인터럽트, 시스템 콜과 같은 트랩을 처리.

User Program

Kernel (Trap Handling)

```
|
|
| -- [System Call] --> | --> Handle syscall (e.g., sys_sigalarm)
|
|
| <---- [Return] -----|
|
|
| -- [Execution] --> Timer Interrupt |
|
|
| ----- [Trap to Kernel] -----|
|
|
```

`usertrap()` checks which interrupt occurred

If timer and alarm condition met:

- Save state
- Set `epc` to alarm handler

```
|
| --- [Return to Alarm Handler] --> User alarm handler executes
|
| -- [sigreturn System Call] --> |
| ----- [Trap to Kernel] -----|
|
```

- Restore saved state
- Resume original code execution

Proc.c - allocproc

```
149 // * labs4 - #3
150 // Set up alarm fields
151 p->alarm_req = 0; // alarm required = 1, not = 0
152 p->alarm_ticks = 0; // alarm unit(period)
153 p->ticks_cnt = 0; // current tick cnt (to check it reaches the top(of tick))
154 p->alarm_handler = 0; // handler address
155 p->alarm_active = 0; // alarm active=1, inactive=0
156 memset(&p->saved_trapframe, 0, sizeof(p->saved_trapframe)); // initialize trapframe
157
```

Sysproc.c - handle syscall

```
97 // * labs4 - #3
98 uint64
99 sys_sigalarm(void)
100 {
101     int ticks;
102     uint64 handler;
103
104     // get args from user - argraw checks for case
105     argint(0, &ticks);
106     argaddr(1, &handler);
107
108     // handle disabling the alarm
109     if(ticks == 0 && handler == 0)
110     {
111         myproc()->alarm_req=0;
112         myproc()->alarm_active=0;
113         myproc()->alarm_ticks=0;
114         myproc()->ticks_cnt=0;
115         myproc()->alarm_handler=0;
116         return 0;
117     }
118
119     // call alarm
120     sigalarm(ticks, handler);
121     return 0;
122 }
123
124 // * labs4 - #3
125 uint64
126 sys_sigreturn(void)
127 {
128     sigreturn();
129     return 0;
130 }
131
```

Proc.h - proc struct

```
108 // * labs4 - #3
109 // alarm fields
110 int alarm_req;
111 int alarm_ticks;
112 int ticks_cnt;
113 uint64 alarm_handler;
114 int alarm_active;
115 struct trapframe saved_trapframe; // copy of registers before interrupt
116 uint64 prev_a0; // value a0 before execution (needed to restore)
117 };
118
```

Proc.c - sig (system call) ...

```
231 // * labs4 - #3
232 int
233 sigalarm(int ticks, uint64 handler)
234 {
235     struct proc *p = myproc();
236
237     p->alarm_req = 1;
238     p->alarm_ticks = ticks;
239     p->ticks_cnt = 0;
240     p->alarm_handler = handler;
241     p->alarm_active = 0;
242     memset(&p->saved_trapframe, 0, sizeof(p->saved_trapframe));
243
244     return 0;
245 }
246
247 // * labs4 - #3
248 int
249 sigreturn(void)
250 {
251     struct proc *p = myproc();
252
253     p->alarm_active = 0;
254     memmove(p->trapframe, &p->saved_trapframe, sizeof(struct trapframe));
255     p->ticks_cnt = 0;
256
257     return p->trapframe->a0;
258 }
259
```

```
149 // * labs4 - #3
150 // Set up alarm fields
151 p->alarm_req = 0; // alarm required = 1, not = 0
152 p->alarm_ticks = 0; // alarm unit(period)
153 p->ticks_cnt = 0; // current tick cnt (to check it reaches the top(of tick))
154 p->alarm_handler = 0; // handler address
155 p->alarm_active = 0; // alarm active=1, inactive=0
156 memset(&p->saved_trapframe, 0, sizeof(p->saved_trapframe)); // initialize trapframe
...
```


Trap.c - usertrap()

```
41 void
42 usertrap(void)
43 {
44     int which_dev = 0;
45
46     if((r_sstatus() & SSTATUS_SPP) != 0)
47         panic("usertrap: not from user mode");
48
49     // send interrupts and exceptions to kerneltrap(),
50     // since we're now in the kernel.
51     w_stvec((uint64)kernelvec);
52
53     struct proc *p = myproc();
54
55     // save user program counter.
56     p->trapframe->epc = r_sepc();
57
58     if(r_scause() == 8){
59         // system call
60
61         if(killed(p))
62             exit(-1);
63
64         // sepc points to the ecall instruction,
65         // but we want to return to the next instruction.
66         p->trapframe->epc += 4;
67
68         // an interrupt will change sepc, scause, and sstatus,
69         // so enable only now that we're done with those registers.
70         intr_on();
71
72         // * labs4 - #3
73         // capture system call number
74         uint64 syscall_num = p->trapframe->a7;
75         syscall();
76
77         // * if sigreturn is called, restore saved a0 value
78         if (syscall_num == SYS_sigreturn)
79         {
80             p->trapframe->a0 = p->saved_trapframe.a0;
81         }
82
83     } else if((which_dev = devintr()) != 0){
84         // ok
85     } else {
86         printf("usertrap(): unexpected scause 0x%lx pid=%d\n", r_scause(), p->pid);
87         printf("             sepc=0x%lx stval=0x%lx\n", r_sepc(), r_stval());
88         setkilled(p);
89     }
90
91     if(killed(p))
92         exit(-1);
93
94     // give up the CPU if this is a timer interrupt.
95     // * labs4 - #3
96     if(which_dev == 2)
97     {
98         if(p->alarm_req && !p->alarm_active)
99         {
100             if (++p->ticks_cnt == p->alarm_ticks)
101             {
102                 p->ticks_cnt = 0;
103                 p->alarm_active = 1;
104                 memmove(&p->saved_trapframe, p->trapframe, sizeof(struct trapframe));
105                 p->trapframe->epc = (uint64)p->alarm_handler; // jump to handler
106             }
107         }
108         yield();
109     }
110     usertrapret();
111 }
```

Trap.c - usertrap()

```
72 // * labs4 - #3
73 // capture system call number which is called by user
74 uint64 syscall_num = p->trapframe->a7;
75 syscall();
76
77 // * if sigreturn is called, restore saved a0 value
78 if (syscall_num == SYS_sigreturn)
79 {
80     p->trapframe->a0 = p->saved_trapframe.a0;
81 }
82
```

Trap.c - usertrap()

```
93
94 // give up the CPU if this is a timer interrupt.
95 // * labs4 - #3
96 if(which_dev == 2)
97 {
98     if(p->alarm_req && !p->alarm_active) // alarm is required && not in execution
99     {
100         if (++p->ticks_cnt == p->alarm_ticks) // increase cnt, and check if it reaches the end of period
101         {
102             p->ticks_cnt = 0; // reset
103             p->alarm_active = 1; // active (prevent re-entrant)
104             memmove(&p->saved_trapframe, p->trapframe, sizeof(struct trapframe)); // save current status
105             p->trapframe->epc = (uint64)p->alarm_handler; // jump to handler
106         }
107     }
108     yield(); // let other process be executed
109 }
110 usertrapret(); // back to user mode (after interrupt or system call)
111 }
112
```

Test case & Key points to implement correctly

- Test0: Basic Alarm Handler Execution
- Test1: Multiple Alarm Handler Calls & Correct State Restoration
- Test2: Reentrant Alarm Prevention
 - `slow_handler()`는 오래 실행되며, 실행 중에 두 번째로 호출되면 안 됨.
- Test3: Register a0 Preservation
 - `sigreturn()`이 호출될 때 레지스터 값이 변조되지 않으면 테스트 통과.

Results

```
enter starting on
$ alarmtest
test0 start
.....alarm!
test0 passed
test1 start
.....alarm!
.....alarm!
.....alarm!
.....alarm!
.....alarm!
.....alarm!
.....alarm!
.....alarm!
.....alarm!
test1 passed
test2 start
.....alarm!
test2 passed
test3 start
test3 passed
^
```

```
Test alarmtest test0 OK (3.8s)
== Test running alarmtest == (3.8s)
== Test alarmtest: test0 ==
alarmtest: test0: OK
== Test alarmtest: test1 ==
alarmtest: test1: OK
== Test alarmtest: test2 ==
alarmtest: test2: OK
== Test alarmtest: test3 ==
alarmtest: test3: OK
== Test usertests == usertests: OK (85.7s)
```