

# Copy on Write

OS Study Session #5

# Some Reviews

# Why OS?

- It's all about ***sharing resources between processes ...***
- Key Requirements
  - Isolation
  - Multiplexing
  - Interaction
- What resources to share?
  - Memory -> This is what we mostly addressed so far
    - Virtual memory, Paging, ...
  - CPU
    - Scheduling ..
  - I/O Device

# Abstraction of Physical Resources

- **Virtual Address**

- Giving an illusion that each processes has its own space
- Should be mapped to valid **Physical Address**

- **Address Translation**


- OS translates VA to PA by its own rule
- Uses Page Table (cf. Translation Lookaside Buffer)

- **Paging / Segmentation**

- VA is broken into fixed-size pages or variable-size segments

# Xv6 Processes and Virtual Address Space

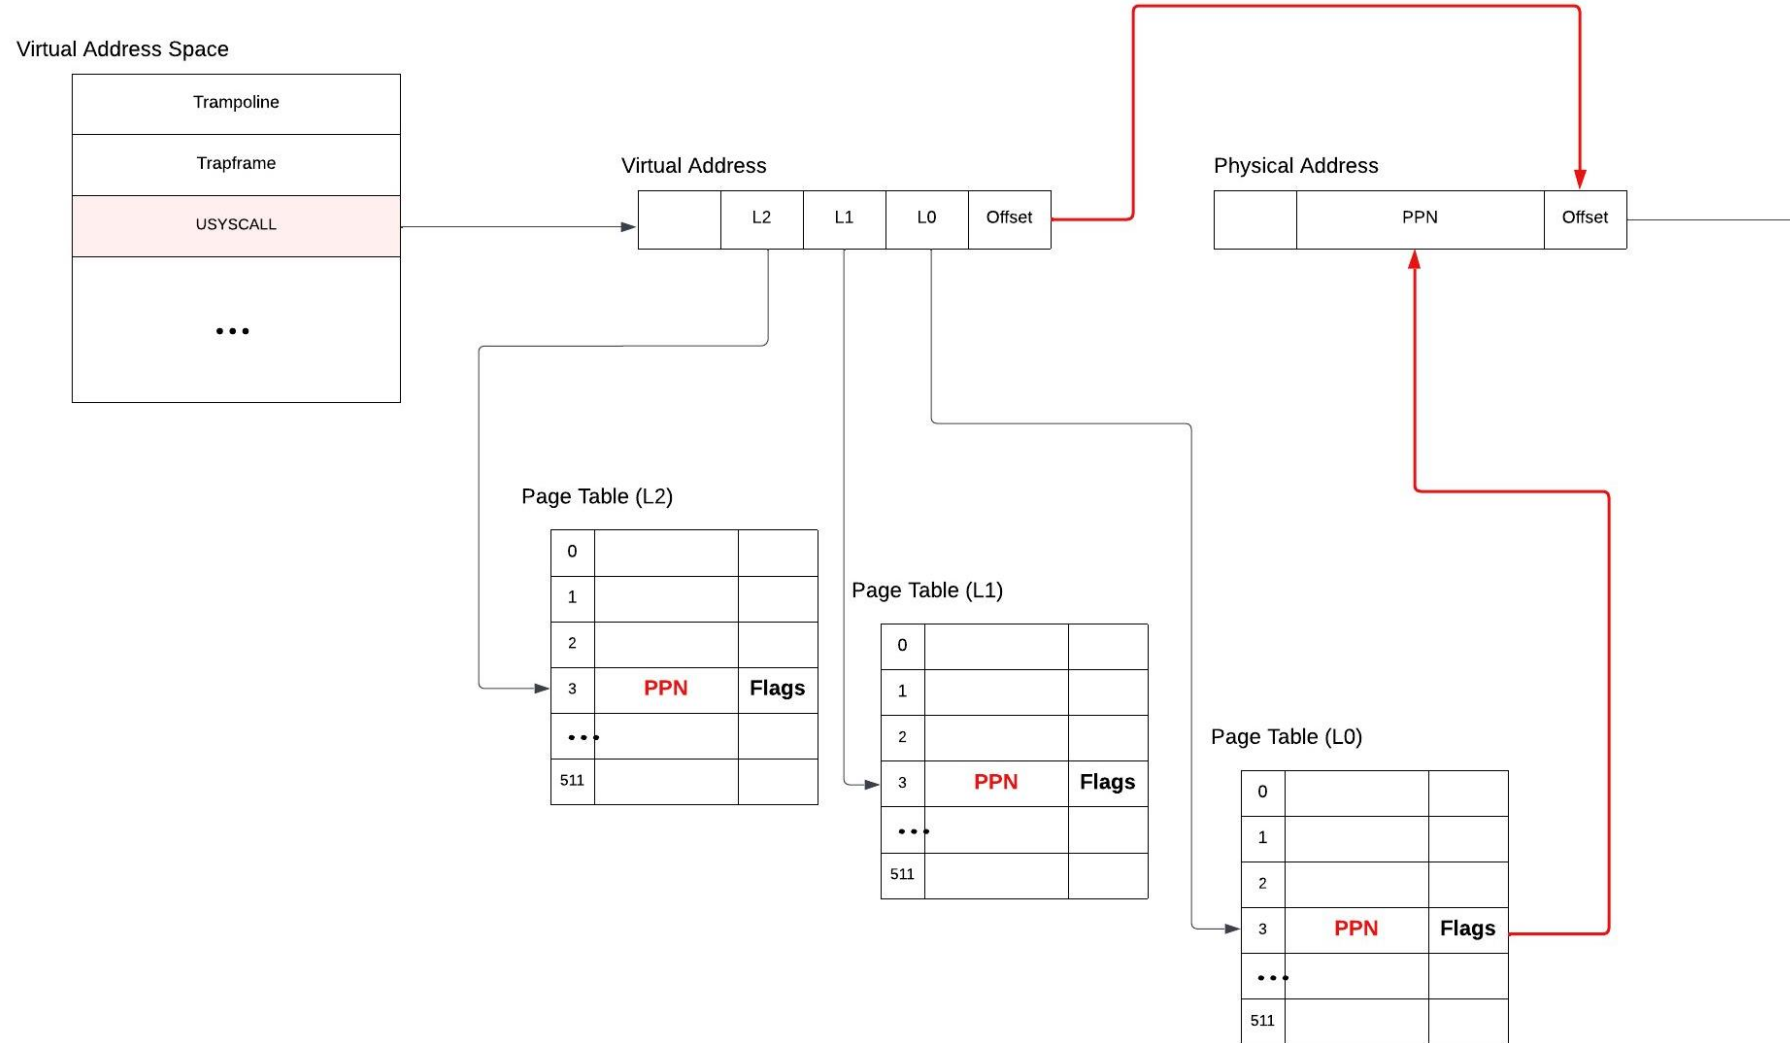
- Process
  - A fundamental unit of job in OS
  - A unit of isolation
- Components
  - Code (Text Section): The program's executable instructions
  - Data Section: Global and static variables
  - Heap: Dynamic memory allocation
  - Stack: Function calls, local variables, return addresses.



Layout of Virtual Address Space

# Page Table Structure

## Full Picture



# Copy-on-Write Concept

# Basic fork() Behavior

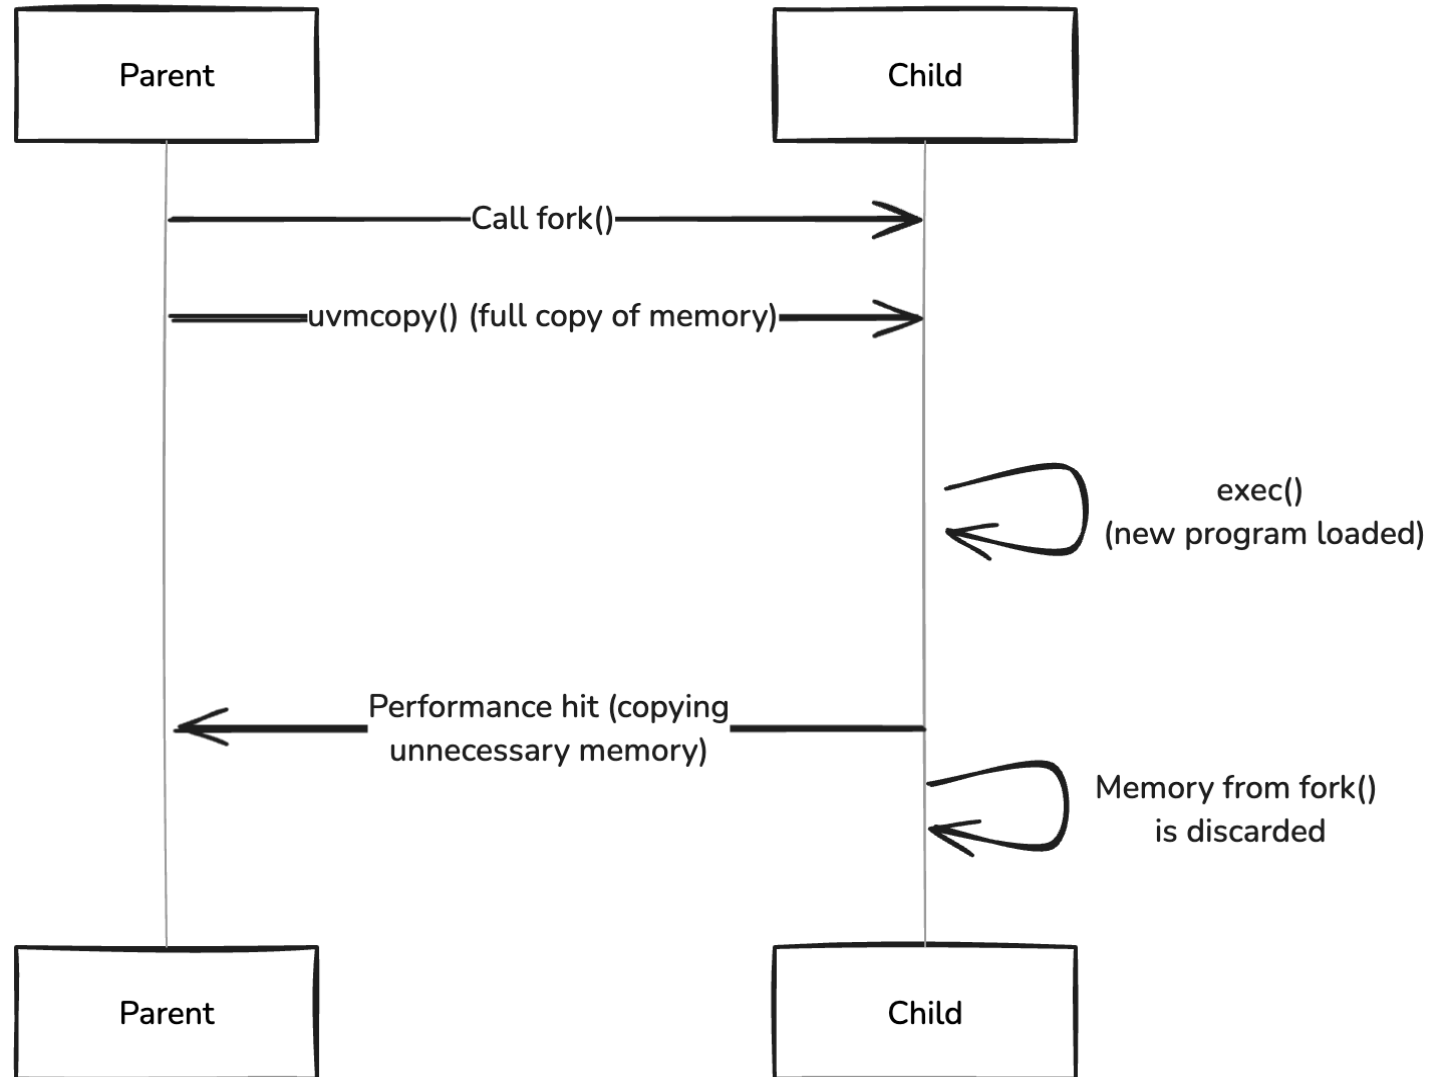
- Fork() creates a new process by duplicating the parent process's memory
- The entire memory of the parent is copied to the child process
- However, if the child immediately calls exec(), all that memory is discarded because the new program loaded by exec() does not use the memory that was copied.



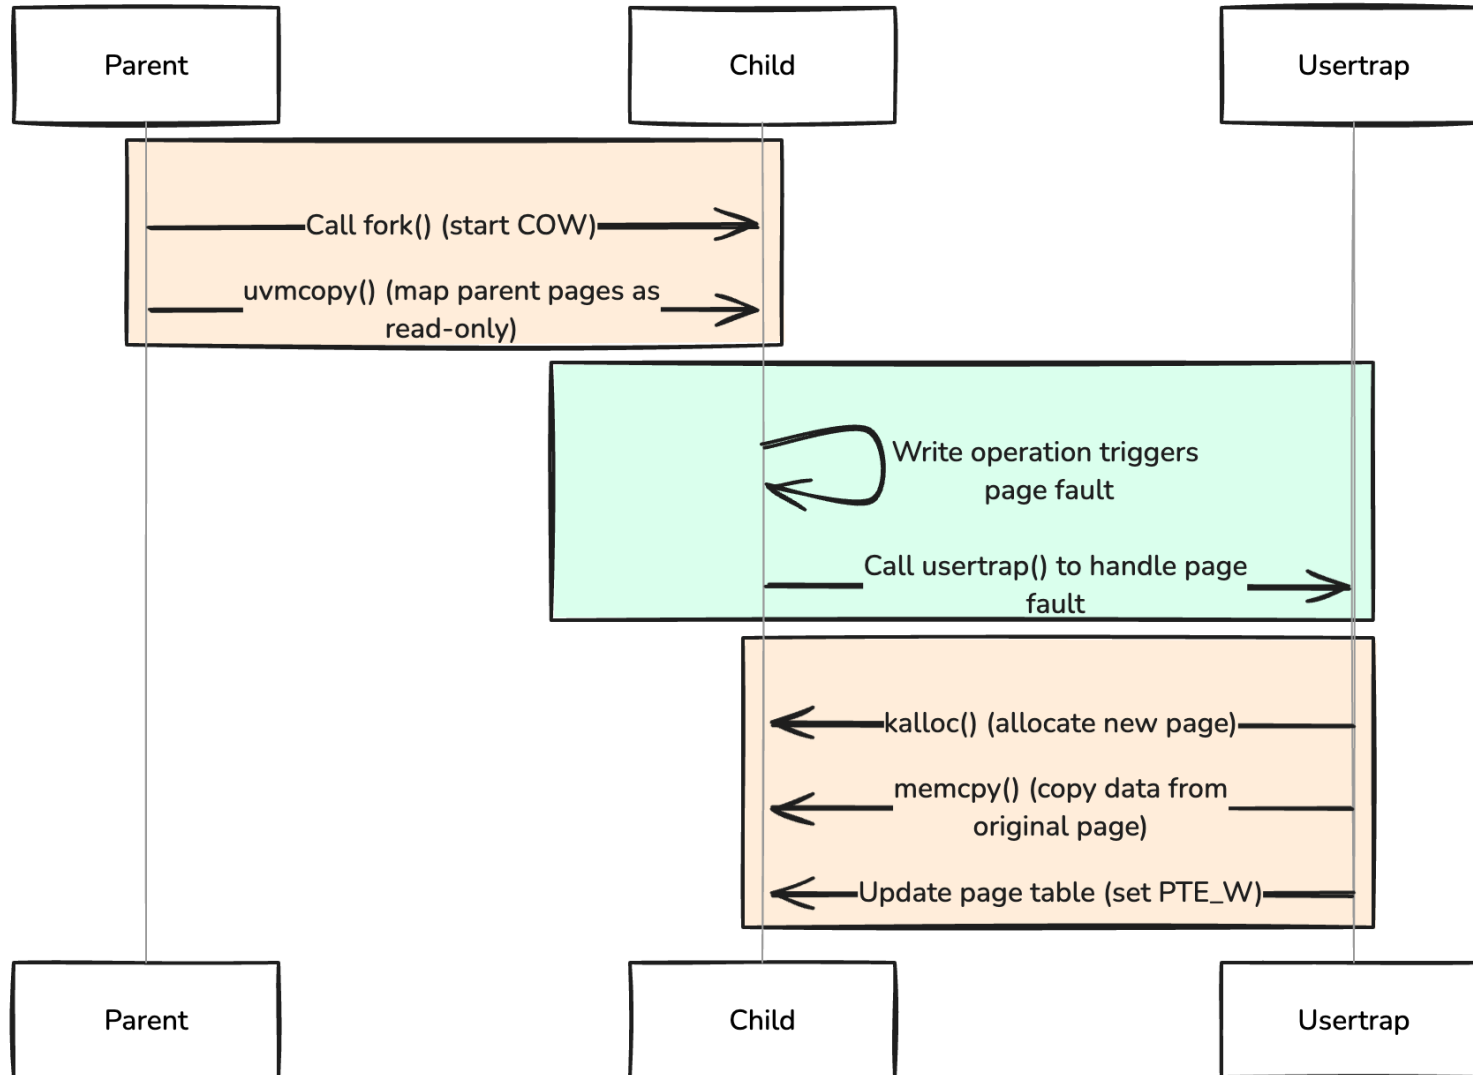
Unnecessary Memory Duplication



# Basic fork() Behavior



# Copy-on-Write Overview




# Key Functions for Implementing COW

- uvmcopy()


[Before COW]

- uvmcopy() performs a full memory copy of the parent process's memory
- All Pages of the parent parent are copied to the child

 Independent copies of memory for parent and child

[After COW]

- uvmcopy() maps parent's memory pages to the child's page table without copying
- Pages are shared, marked as read-only

 Pages are only copied when either process attempts to write.

# Key Functions for Implementing COW

- usertrap()

[Before COW]

- usertrap() handles page faults caused by invalid memory access (e.g. accessing unmapped memory)

[After COW]

- usertrap() handles COW page faults and allocates a new page only when a write occurs on a shared, read-only page

# Key Functions for Implementing COW

- copyout()

[Before COW]

- copyout() simply copies data from kernel space to user space

[After COW]

- copyout() ensures that data from a COW page is correctly transeferred from kernel space to user space

# Reference counting for Implementing COW

## [Before COW]

- No reference counting is implemented, meaning pages are freed immediately after a process finishes using them

## [After COW]

- Reference counting is introduced to track how many processes are using a page.
- `kalloc()` and `kfree()` are modified to support reference counting

Labs:

Implement copy-on-write fork

(hard)

# Problem Summary

- the `fork()` system call currently copies all of the parent process's memory into the child process - inefficient
- Inefficient when `fork()` is followed by `exec()` which discards most of that copied memory



# Implementation Tasks

1. Instead of copying all memory during fork, just create a page table for the child with PTEs pointing to the parent's physical pages
2. Mark all user PTEs in both parent and child as read-only
3. When either process tries to write to a page, it causes a page fault
4. The kernel handler then:
  - Allocates a new physical page
  - Copies the content from the original page
  - Updates the PTE to point to the new page with write permission
5. Implement reference counting for physical pages to free them only when no process references them

# Implementation Tasks

1. Use RSW bits in RISC-V PTE to mark COW pages
2. Modify `uvmcopy()` to map parent's pages without copying and mark pages as read-only
3. Modify `usertrap()` to handle COW page faults
4. Implement reference counting for physical pages in `kalloc.c`
5. Modify `copyout()` to handle COW pages
6. Modify `kfree()` to free page only when no process references them

## 1. Use RSW bits in RISC-V PTE to mark COW pages

```
#define PTE_FLAGS(pte) ((pte) & 0x3FF)  
+ #define PTE_COW (1L << 8) // Copy-on-write page
```

- Add COW bit to PTE (riscv.h)
- Flag for marking if this page is COW or not
- Enables recognition between COW page fault and normal faults

```

int
uvmcopy(pagetable_t old, pagetable_t new, uint64 sz)
{
    pte_t *pte;
    uint64 pa, i;
    uint flags;

    for(i = 0; i < sz; i += PGSIZE){
        if((pte = walk(old, i, 0)) == 0)
            panic("uvmcopy: pte should exist");
        if((*pte & PTE_V) == 0)
            panic("uvmcopy: page not present");
        pa = PTE2PA(*pte);
        flags = PTE_FLAGS(*pte);

        // If the page is writeable, mark it COW
        if(flags & PTE_W) {
            flags &= ~PTE_W; // Clear writable bit
            flags |= PTE_COW; // Set COW bit
            *pte = PA2PTE(pa) | flags;
        }

        // Map the page in child's page table
        if(mappages(new, i, PGSIZE, pa, flags) != 0)
            goto err;

        incref((void*)pa);
    }
    return 0;

err:
    uvmunmap(new, 0, i / PGSIZE, 1);
    return -1;
}

```

vm.c

## 2. Modify uvmcopy ( ) to map parent's pages without copying and mark pages as read-only

Change the writable page to read-only  
And set the COW bit

Point to same physical page  
instead of copying

Increment page reference count

### 3. Modify usertrap() to handle COW page faults

```
syscall();
} else if((which_dev = devintr()) != 0){
    // ok
} else if(r_scause() == 15) { // Page fault on store
    uint64 va = r_stval();
    if(handle_cow_fault(p->pagetable, va) < 0) {
        printf("usertrap(): cow page fault failed\n");
        p->killed = 1;
    }
} else {
    printf("usertrap(): unexpected scause %lx pid=%d\n", r_scause(), p->pid);
    printf("          sepc=%lx stval=%lx\n", r_sepc(), r_stval());
    p->killed = 1;
}

if(killed(p))
    exit(-1);

// give up the CPU if this is a timer interrupt.
if(which_dev == 2)
    yield();

usertrapret();
}
```

trap.c

If the page fault is store page fault (scause==15)

1. Check if the fault is from COW page
2. Call handle\_cow\_fault()

```

int
handle_cow_fault(pagetable_t pagetable, uint64 va)
{
    pte_t *pte;
    uint64 pa;
    uint flags;
    char *mem;

    if(va >= MAXVA) return -1;

    va = PGROUNDDOWN(va);
    pte = walk(pagetable, va, 0);

    if(pte == 0) return -1;
    if((*pte & PTE_V) == 0 || (*pte & PTE_U) == 0) return -1;
    if((*pte & PTE_COW) == 0) return -1;

    pa = PTE2PA(*pte);
    flags = PTE_FLAGS(*pte);

    // Allocate new page
    if((mem = kalloc()) == 0) return -1;

    // Copy old page to new page
    memmove(mem, (char*)pa, PGSIZE);

    // Map new page with write permission
    flags &= ~PTE_COW; // Clear COW bit
    flags |= PTE_W;    // Set writable bit

    uvmunmap(pagetable, va, 1, 0);
    if(mappages(pagetable, va, PGSIZE, (uint64)mem, flags) != 0) {
        kfree(mem);
        return -1;
    }

    kfree((void*)pa);
    return 0;
}

```

- Check if va is valid
- Find PTE of the address in pagetable

- Check if PTE exists
- Check if PTE is valid & accessible from user mode
- Check if this page is actually COW

- Extract physical address, flag bits from PTE

- Allocate new physical memory
  - Fails if not enough memory
- Copy old page to new page

- Not a COW page anymore (writable)

- Remove old mapping
  - do\_free=0 : don't release physical memory
- Map new page to va with updated flags
  - Release allocated memory and return error if mapping fails

## 4. Implement reference counting for physical pages in `kalloc.c`

```
+ // Reference count for each physical page
+ #define PA2IDX(pa) (((uint64)pa - KERNBASE) / PGSIZE)
+ #define MAX_PAGES ((PHYSTOP - KERNBASE) / PGSIZE)
+ struct { // Count reference for a page
+     struct spinlock lock;
+     int count[MAX_PAGES];
+ } ref_count;
```

```
void
kinit()
{
    initlock(&kmem.lock, "kmem");
+   initlock(&ref_count.lock, "ref_count");
    freerange(end, (void*)PHYSTOP);
}
```

```
+ // Increment reference count for a page
+ void
+ incref(void *pa)
+ {
+     acquire(&ref_count.lock);
+     ref_count.count[PA2IDX(pa)]++;
+     release(&ref_count.lock);
+ }
+
+ // Get reference count for a page
+ int
+ getref(void *pa)
+ {
+     int count;
+     acquire(&ref_count.lock);
+     count = ref_count.count[PA2IDX(pa)];
+     release(&ref_count.lock);
+     return count;
+ }
+
```

```

int
copyout(pagetable_t pagetable, uint64 dstva, char *src, uint64 len)
{
    uint64 n, va0, pa0;
    pte_t *pte;

    while(len > 0){
        va0 = PGROUNDDOWN(dstva);
        if(va0 >= MAXVA)
            return -1;

        pte = walk(pagetable, va0, 0);
        if(pte == 0 || (*pte & PTE_V) == 0 || (*pte & PTE_U) == 0)
            return -1;

        // Handle COW page
        if((*pte & PTE_W) == 0 && (*pte & PTE_COW)) {
            if(handle_cow_fault(pagetable, va0) < 0)
                return -1;
            pte = walk(pagetable, va0, 0);
        }

        if((*pte & PTE_W) == 0)
            return -1;

        pa0 = PTE2PA(*pte);
        n = PGSIZE - (dstva - va0);
        if(n > len)
            n = len;
        memmove((void *) (pa0 + (dstva - va0)), src, n);

        len -= n;
        src += n;
        dstva = va0 + PGSIZE;
    }
    return 0;
}

```

## 5. Modify copyout() to handle COW pages

Copy data from kernel space to user space.  
Used in file-read or returning syscall result

- Find va in page table
- Return -1 if No PTE, invalid PTE, not user-accessible

If read-only COW page

- Call handle\_cow\_fault()
- Find the updated PTE pointing to a new page

- Get physical address from PTE
- Copy data



## 6. Modify kfree( ) to free page only when no process references them

```
void
kfree(void *pa)
{
    struct run *r;

    if(((uint64)pa % PGSIZE) != 0 || (char*)pa < end || (uint64)pa >= PHYSTOP)
        panic("kfree");

    acquire(&ref_count.lock);
    if(--ref_count.count[PA2IDX(pa)] > 0) { // if (cnt == 0) release mem
        release(&ref_count.lock);
        return;
    }
    release(&ref_count.lock);

    // Fill with junk to catch dangling refs.
    memset(pa, 1, PGSIZE);

    r = (struct run*)pa;

    acquire(&kmem.lock);
    r->next = kmem.freelist;
    kmem.freelist = r;
    release(&kmem.lock);
}
```

Release page if no process is referencing the page

- Cast to-be-released page into freelist struct
- Add released page to freelist

