

# Network Driver

OS Study Session #6

Labs

Part One: NIC

## Problem Summary

- Filling the function `e1000_recv()` and `e1000_transmit()`
- Understanding overall process of sending / receiving data with NIC

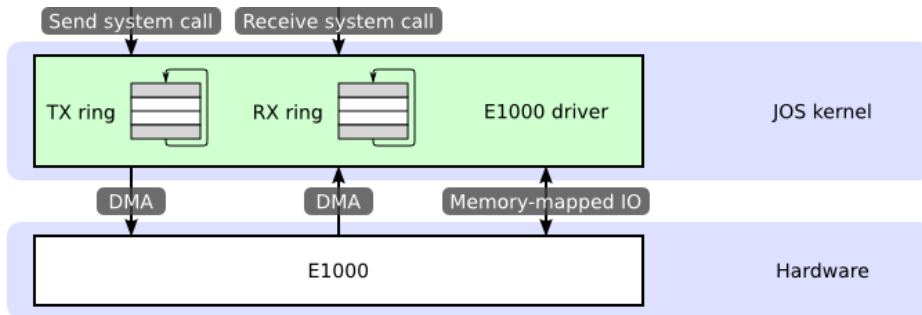
## Role of Network Driver (E1000)

- Allocating memory for the transmit and receive queues
- Setting up DMA descriptors
- Configuring the E1000 with the location of these queues
- everything after that is asynchronous

DMA (Direct Memory Access) : read and write packet data directly from memory without involving the CPU


# Receive System Call

When the E1000 receives a packet, it copies it into the next DMA descriptor in the receive queue, which the driver can read from at its next opportunity.



1. NIC receives the new data packet in receiving buffer
2. An interrupt is triggered by NIC to notify
3. Network driver calls `e1000_rcv()` to handling this interrupt (`e1000_rcv()` = interrupt handler)
4. `e1000_rcv()` transfers the data to OS's network stack

# Receive Descriptor Ring Structure



- The driver always add descriptors to the tail and moves the tail pointer
- Hardware always consumes descriptors from the head and moves the head pointer
- The contents of the queue are the descriptors between head and tail pointer
- Receive queue contents : Free descriptors that the card can receive packets into
- Transmit queue content : packets waiting to be sent

# e1000\_recv

```
#define E1000_RXD_STAT_DD 0x01 /* Descriptor Done */  
#define E1000_RXD_STAT_EOP 0x02 /* End of Packet */
```

net\_rx() : deliver a packet to the networking stack

1. Find an the next waiting received packet index
2. Check if a new packet is available
3. Deliver a packet
4. Allocate a new buffer (kalloc()) to replace the one just given to net\_rx().
5. Update E1000\_RDT register to the new index

```
static void  
e1000_recv(void)  
{  
    //  
    // Your code here.  
    //  
    // Check for packets that have arrived from the e1000  
    // Create and deliver a buf for each packet (using net_rx()).  
    //  
  
    while (1) {  
        int idx = (regs[E1000_RDT] + 1) % RX_RING_SIZE;  
  
        if ((rx_ring[idx].status & E1000_RXD_STAT_DD) == 0) {  
            return;  
        }  
  
        if (rx_ring[idx].status & E1000_RXD_STAT_EOP) {  
  
            int len = rx_ring[idx].length;  
  
            net_rx(rx_bufs[idx], len);  
            rx_bufs[idx] = kalloc();  
            if (rx_bufs[idx] == 0) {  
                return;  
            }  
            rx_ring[idx].status = 0;  
            rx_ring[idx].addr = (uint64_t)rx_bufs[idx];  
        }  
        regs[E1000_RDT] = idx;  
    }  
}
```

# e1000\_transmit

```
#define E1000_TXD_CMD_RS 0x08000000 /* Report Status */  
#define E1000_TXD_CMD_EOP 0x01000000 /* End of Packet */
```

To transmit a packet, the driver copies it into the next DMA descriptor in the transmit queue and informs the E1000 that another packet is available;

[https://pdos.csail.mit.edu/6.1810/2024/readings/8254x\\_GBe\\_SDM.pdf](https://pdos.csail.mit.edu/6.1810/2024/readings/8254x_GBe_SDM.pdf)

1. Finding the next packet index  
(by E1000\_TDT control register)
2. Check overFlowed / not completed
3. Free the last buffer (kfree()) if there was
4. Fill the descriptor

```
int  
e1000_transmit(char *buf, int len)  
{  
    // buf contains an ethernet frame; program it into  
    // the TX descriptor ring so that the e1000 sends it. Stash  
    // a pointer so that it can be freed after send completes.  
    //  
    // Acquire lock  
    acquire(&e1000_lock);  
  
    // get tail index  
    int idx = regs[E1000_TDT];  
  
    // check whether it is overflowed/not completed  
    if ((tx_ring[idx].status & E1000_TXD_STAT_DD) == 0) {  
        release(&e1000_lock);  
        return -1;  
    }  
  
    // free the last buffer  
    if (tx_bufs[idx])  
        kfree(tx_bufs[idx]);  
  
    tx_bufs[idx] = buf;  
    tx_ring[idx].length = len;  
    tx_ring[idx].addr = (uint64) (buf);  
    tx_ring[idx].cmd = E1000_TXD_CMD_RS | E1000_TXD_CMD_EOP;  
    tx_ring[idx].status = 0;  
  
    // update ring location  
    regs[E1000_TDT] = (idx + 1) % TX_RING_SIZE;  
  
    release(&e1000_lock);  
    return 0;  
}
```



Labs

Part Two: UDP Receive

# Problem Summary

- A system call for **sending UDP packets** has already been implemented.
- The objective is to create a system call for **receiving UDP packets**

# Related system-call API and function

- **send(sport, dst, dport, \*buf, len)** (→ Already implemented)
  - Sends a packet to the destination specified by (dst, dport).
  - **Packet:** Starts at buf and has a length of len.
- **recv(dport, \*src, \*sport, \*buf, len)** (→ To be implemented)
  - Receives packets arriving at dport.
  - Writes the received information into src and sport, which are virtual addresses.
- **bind** (→ To be implemented)
  - A process must bind to a port before calling recv.
- **ip\_rx** (→ To be implemented)
  - Restores the packet before calling recv.

ports

0

1

2

...

p

...

max port #

Packet arrived to unbound port -> discarded



pkt0

1. bind



struct

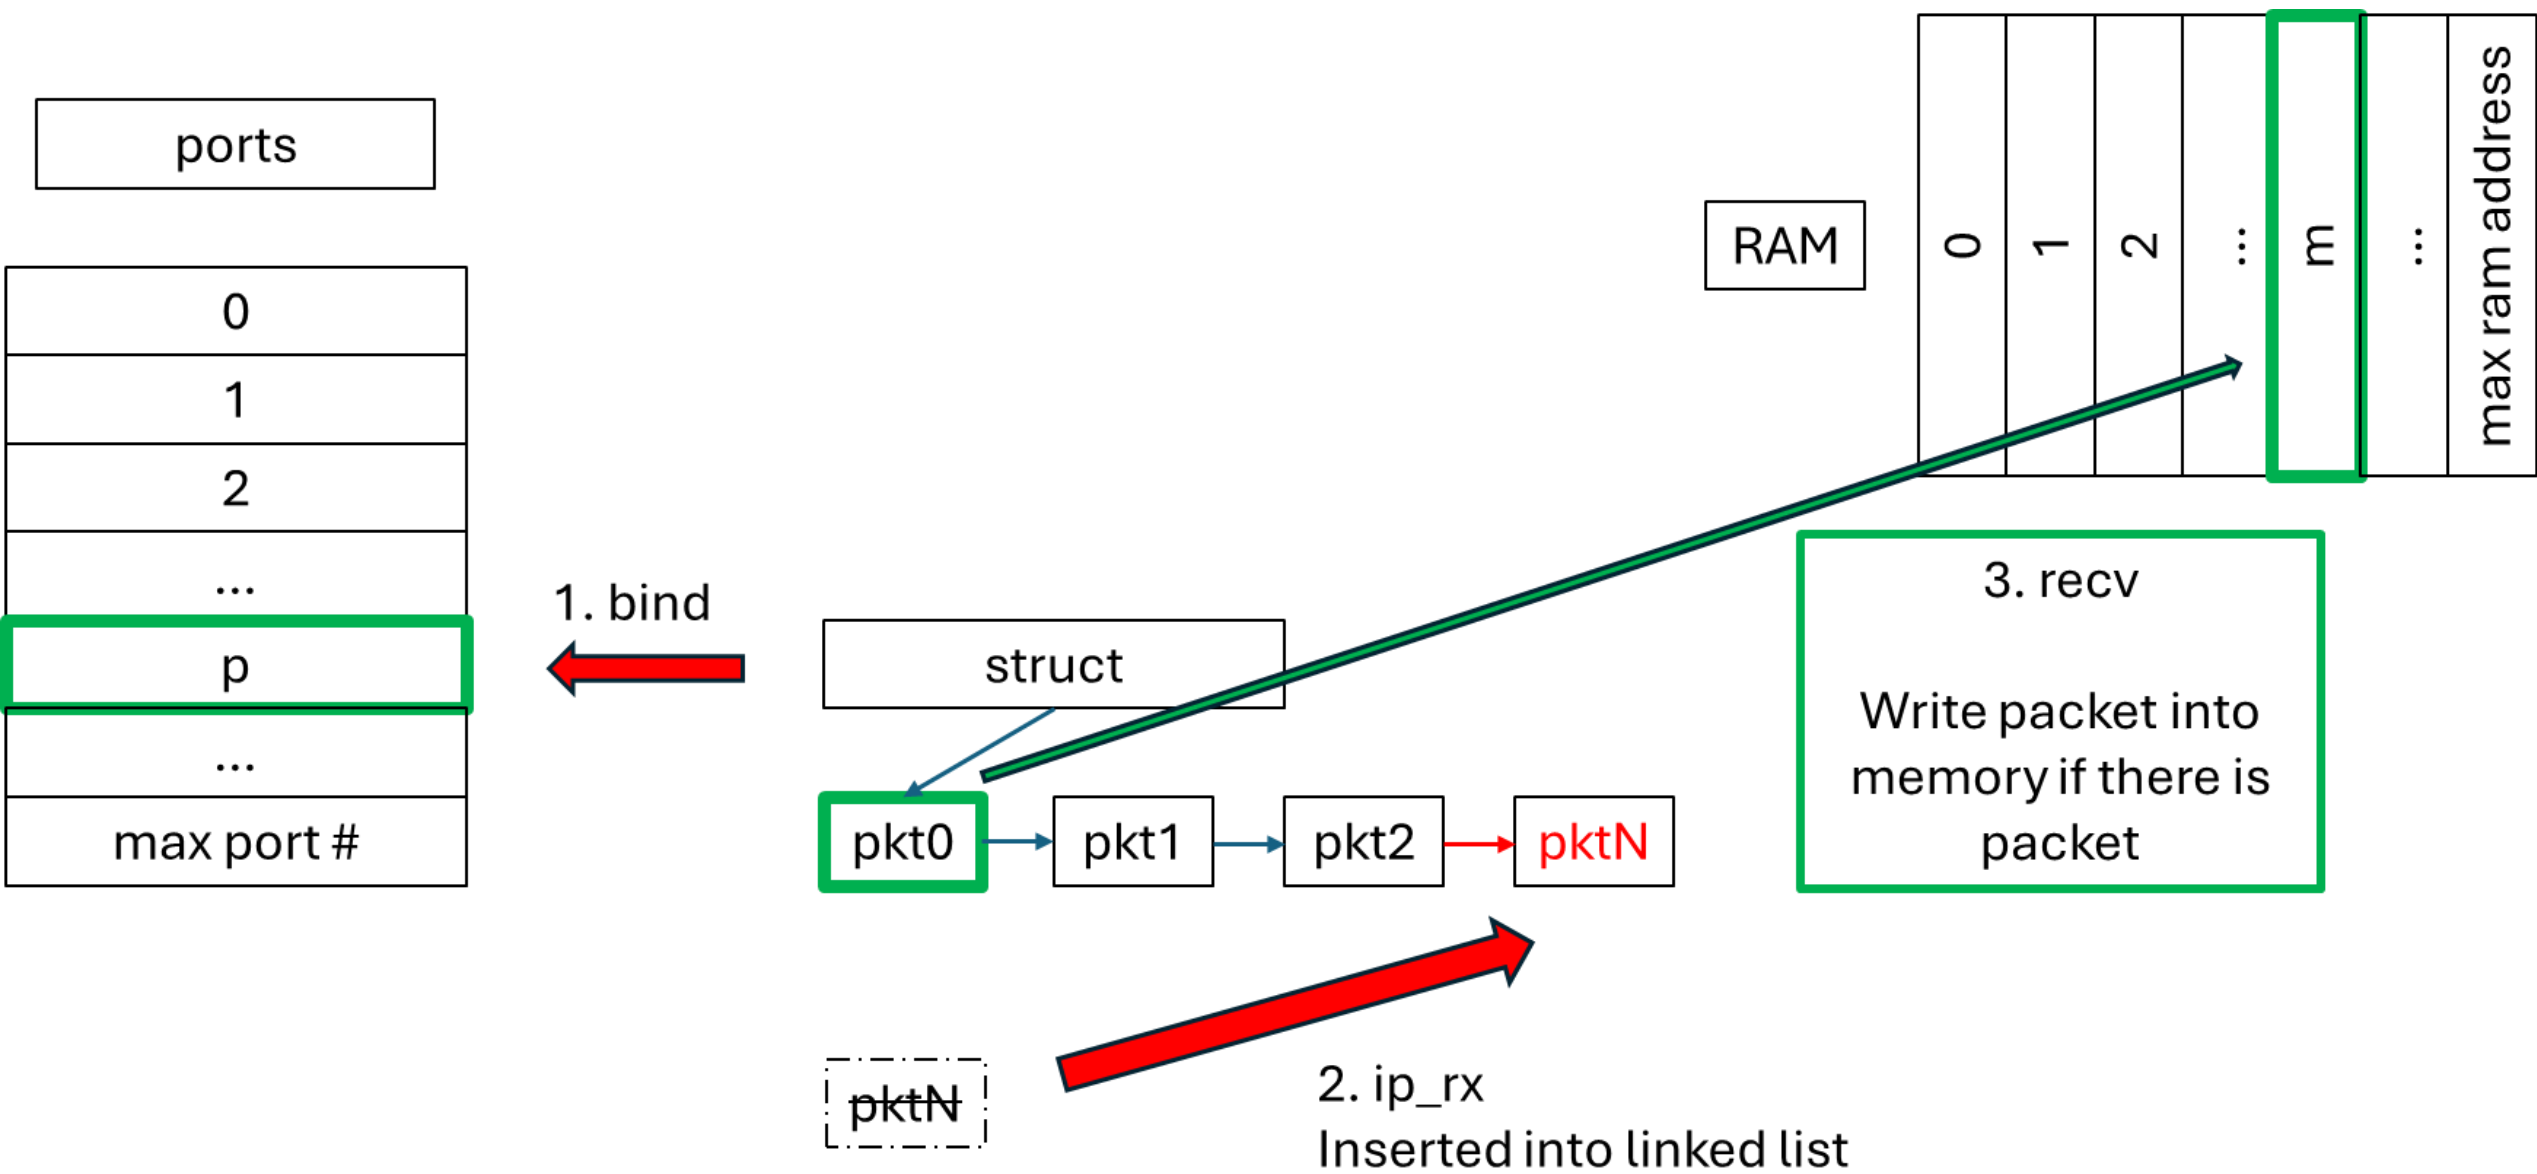
pkt0

pkt1

pkt2

Packet arrived to bound port

pktN



# bind(port)

- Binding **port** before calling **ip\_rx** and **recv**
- The maximum number of packets that can arrive is limited to 16
- If a packet arrives at an unbound port, the packet should be discarded

# Data structure

- Create "struct **bound\_port**"
- Each bound port must store waiting packets.
- A list of **struct bound\_port** instances is stored in the **udp\_table**.
- The function **find\_bound\_port** is called to check whether the given port is bound or not.

```
struct bound_port
{
    uint16 port;          // Bound port number
    int pending_count;    // Number of pending packets
    struct spinlock lock;
    struct pending_packet *head; // Packet queue head
    struct pending_packet *tail; // Packet queue tail
};

#define MAX_PORTS 64
struct
{
    struct spinlock lock;
    struct bound_port ports[MAX_PORTS];
    int port_count;
} udp_table;

void
netinit(void)
{
    initlock(&netlock, "netlock"); // Initialize lock
    udp_table.port_count = 0;        // Initialize port counter
}
```

```
static struct bound_port *
find_bound_port(uint16 port)
{
    for (int i = 0; i < udp_table.port_count; i++)
    {
        if (udp_table.ports[i].port == port)
            return &udp_table.ports[i]; // Return pointer to matching port structure
    }
    return 0; // Return 0 if not found
}
```

# bind(port)

- A lock is necessary since other processes may try to bind the same port.
- If the port is already bound or the number of bound ports has already exceeded **MAX\_PORTS**, the binding fails.

```
uint64
sys_bind(void)
{
    int port; argint(0, &port);

    acquire(&udp_table.lock);

    {
        // Check if the port is already bound
        if (find_bound_port(port) != 0 || udp_table.port_count >= MAX_PORTS)
        {
            release(&udp_table.lock);
            return -1;
        }

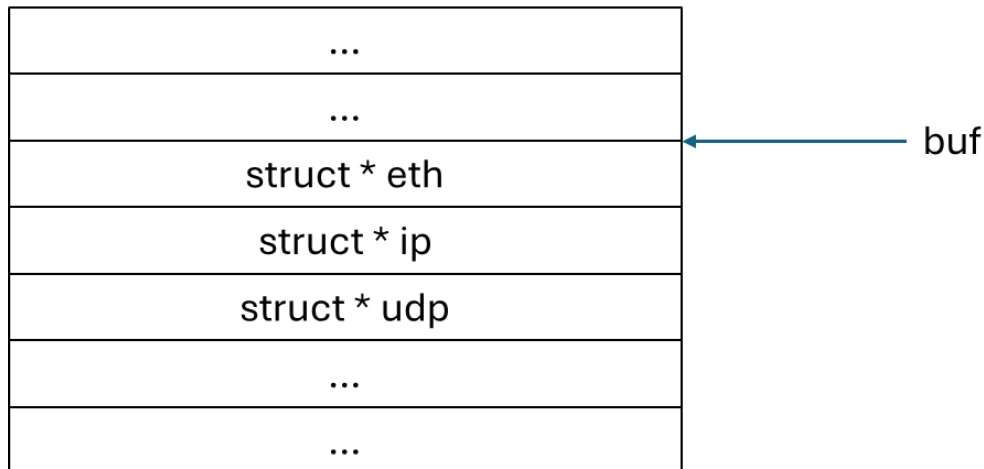
        // Initialize a new bound_port structure
        struct bound_port *bp = &udp_table.ports[udp_table.port_count++];
        bp->port = port;
        bp->pending_count = 0;
        bp->head = bp->tail = 0;
        initlock(&bp->lock, "bound_port");
    }

    release(&udp_table.lock);
    return 0;
}
```



# ip\_rx(buf, len)

- This function must receive the packet and store it in the queue within bound\_port.
- Starting from the buf address, the addresses of the **Ethernet packet**, **IP packet**, and **UDP packet** are stored.
- Packets are represented as structures in net.h.



```
// a UDP packet header (comes after an IP header).
struct udp {
    uint16 sport; // source port
    uint16 dport; // destination port
    uint16 ulen;  // length, including udp header, not including IP header
    uint16 sum;   // checksum
};
```

# Step1 of ip\_rx(buf, len)

- Get the destination port from the UDP header.
- The function ntohs rearranges the byte order.
  - Packets use Big-endian
  - CPU use Little-endian


```
void
ip_rx(char *buf, int len)
{
    // don't delete this printf; make grade depends on it.
    static int seen_ip = 0;
    if(seen_ip == 0)
        printf("ip_rx: received an IP packet\n");
    seen_ip = 1;

    struct eth *ethhdr = (struct eth *)buf;
    struct ip *iphdr = (struct ip *)(ethhdr + 1);
    struct udp *udphdr = (struct udp *)(iphdr + 1);
    uint16 dport = ntohs(udphdr->dport);

    acquire(&udp_table.lock);

    struct bound_port *bp = find_bound_port(dport);
    if (bp == 0)
    {
        kfree(buf);
        release(&udp_table.lock);
        return;
    }

    acquire(&bp->lock);
    release(&udp_table.lock);
}
```



```
struct udp {
    uint16 sport; // source port
    uint16 dport; // destination port
    uint16 ulen;  // length, including udp header, not including IP header
    uint16 sum;   // checksum
};
```

# Step2 of ip\_rx(buf, len)

- If the number of waiting packets is equal to **MAX\_PENDING\_PACKETS**, return from the function.
- If not, initialize the packet.

```
if (bp->pending_count >= MAX_PENDING_PACKETS)
{
    kfree(buf);
    release(&bp->lock);
    return;
}

int payload_len = ntohs(udphdr->ulen) - sizeof(struct udp);

char *data = kalloc();
struct pending_packet *pp = kalloc();

memmove(data, (char *)(udphdr + 1), payload_len);

memset(pp, 0, sizeof(*pp));
pp->data = data;
pp->len = payload_len;
pp->src_ip = ntohl(iphdr->ip_src);
pp->src_port = ntohs(udphdr->sport);
pp->next = 0;
```

# Step3 of ip\_rx(buf, len)

- Insert the prepared packet into the queue.
- It has to call the wakeup function.
  - I will explain later.

```
// Update the queue
if (bp->tail == 0)
{
    bp->head = bp->tail = pp;
}
else
{
    bp->tail->next = pp;
    bp->tail = pp;
}
bp->pending_count++;

// Wake up the waiting process
wakeup(bp);
kfree(buf);
release(&bp->lock);
```

# recv(dport, \*src, \*sport, \*buf, len)

- Extract a packet stored in the queue.
- If no packet has arrived, wait until a new packet arrives.
- Store the IP address of the packet in src and store the port of the packet in sport.

# Step1 of recv(dport, \*src, \*sport, \*buf, len)

- Wait until the packet is received.
- If no packet has arrived, wait (by using sleep).
- It can be awakened by the previously mentioned wakeup.
  - The wakeup function is called if a packet arrives.
  - Exit the while loop and proceed to the next step.

```
acquire(&udp_table.lock);
if ((bp = find_bound_port(dport)) == 0)
{
    release(&udp_table.lock);
    return -1;
}

acquire(&bp->lock);
release(&udp_table.lock);

while (bp->head == 0)
{
    if (p->killed)
    {
        release(&bp->lock);
        return -1;
    }
    sleep(bp, &bp->lock);
}
```

# Step2 of recv(dport, \*src, \*sport, \*buf, len)

- Src and sport are virtual addresses, which will store the source IP and source port, respectively.
- Extract the stored packet from the queue.
- Store the source IP and source port in the given virtual addresses.
  - Using copyout api

```
pp = bp->head;
bp->head = pp->next;
if (bp->head == 0)
{
    bp->tail = 0;
}
bp->pending_count--;

int copy_len = pp->len;
if (copy_len > maxlen)
    copy_len = maxlen;

if (copyout(p->pagetable, src_arg, (char *)&pp->src_ip, sizeof(pp->src_ip)) < 0 ||
    copyout(p->pagetable, sport_arg, (char *)&pp->src_port, sizeof(pp->src_port)) < 0)
{
    goto bad;
}

if (copyout(p->pagetable, buf, pp->data, copy_len) < 0)
{
    goto bad;
}
```